

# CV

Practical No	Practical Name	Page No
Practical-1	Program for reading, writing and displaying images	1
Practical-2	Program for changing color spaces	1
Practical-3	Program for resizing images	2
Practical-4	Program to rotate image	3
Practical-5	Programs using Histogram Equalization	4
Practical-6	Program for Edge Detection	5
Practical-7	Program for Line Detection	6
Practical-8	Programs using Scale Invariant Feature Transform (SIFT)	6
Practical-9	Programs for Motion Detection	11
Practical-10	Programs for Face Detection	122
Practical-11	Object Detection	126

```
#Practical-1-Program for reading, writing and displaying images

import cv2
from matplotlib import pyplot as plt
image = cv2.imread('cr7.jpg')
image_rgb = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
plt.imshow(image_rgb)
plt.axis('off')
plt.show()
output_path = 'image1.jpg'
cv2.imwrite(output_path, image)
print(f"Image written to {output_path}")
```



Image written to image1.jpg

#Practical-2-Program for changing color spaces

```
import cv2
from google.colab.patches import cv2_imshow

# Load your image
path = 'cr7.jpg' # Replace with the path to your image
src = cv2.imread(path)

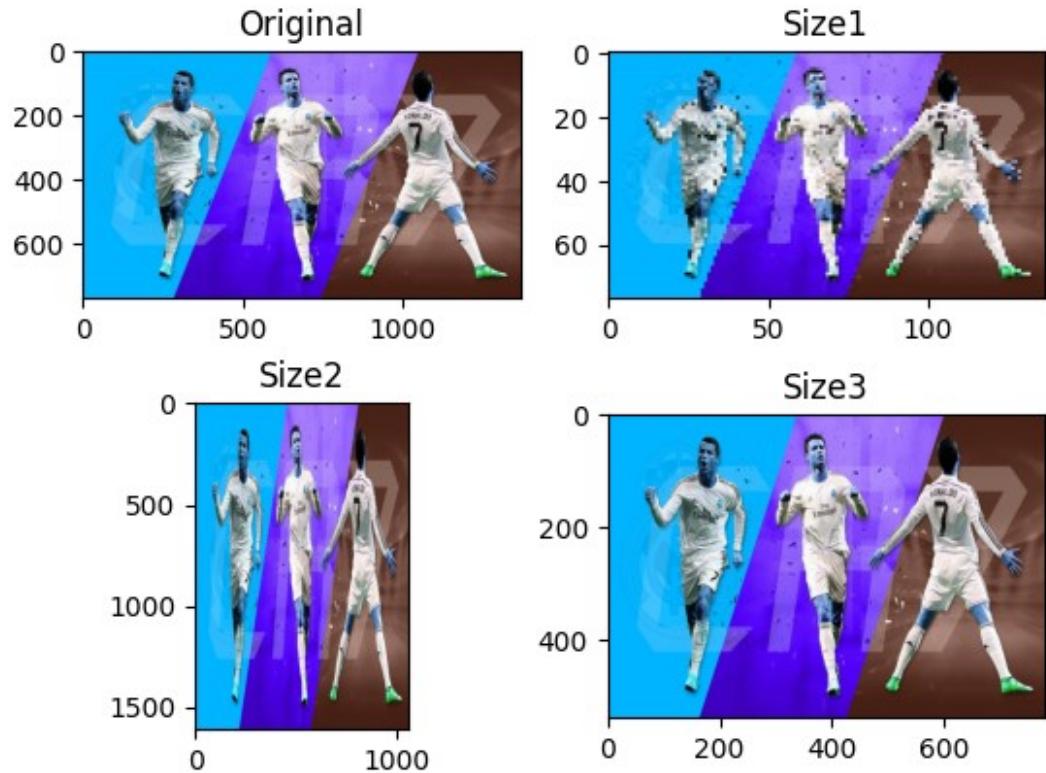
# Convert the image to grayscale
image = cv2.cvtColor(src, cv2.COLOR_BGR2GRAY)

# Display the image using cv2_imshow
cv2_imshow(image)
```



#Practical-3-Program for resizing images

```
import cv2
import numpy as np
import matplotlib.pyplot as plt
image = cv2.imread(r'cr7.jpg')
# Loading the image
size1 = cv2.resize(image, (0, 0), fx = 0.1, fy = 0.1)
size2 = cv2.resize(image, (1050, 1610))
stretch_near = cv2.resize(image, (780, 540), interpolation =
cv2.INTER_LINEAR)
Titles =[ "Original", "Size1", "Size2", "Size3"]
images =[image, size1, size2, stretch_near]
count = 4
for i in range(count):
    plt.subplot(2, 2, i + 1)
    plt.title(Titles[i])
    plt.imshow(images[i])
plt.show()
```



#Practical-4-Program to rotate image

```

import cv2
from matplotlib import pyplot as plt

image = cv2.imread('cr7.jpg')
num_rows, num_cols = image.shape[:2]
rotation_matrix = cv2.getRotationMatrix2D((num_cols / 2, num_rows / 2), 180, 0.6)
img_rotation = cv2.warpAffine(image, rotation_matrix, (num_cols, num_rows))
img_rotation_rgb = cv2.cvtColor(img_rotation, cv2.COLOR_BGR2RGB)
plt.imshow(img_rotation_rgb)
plt.axis('off') # Hide axis plt.show()

(-0.5, 1365.5, 767.5, -0.5)

```



#Practical-5-Programs using Histogram Equalization

```
import numpy as np
import cv2 as cv
from matplotlib import pyplot as plt
img = cv.imread('cr7.jpg', cv.IMREAD_GRAYSCALE)
equ = cv.equalizeHist(img)
res = np.hstack((img, equ))
cv.imwrite('res.png', res)
plt.figure(figsize=(10, 5))
plt.imshow(res, cmap='gray')
plt.axis('off')
plt.show()
print('Image written to res.png')
```



Image written to res.png

```
import cv2 as cv
from matplotlib import pyplot as plt
img = cv.imread('cr7.jpg', cv.IMREAD_GRAYSCALE)
```

```

clahe = cv.createCLAHE(clipLimit=2.0, tileGridSize=(8,8))
cl1 = clahe.apply(img)
cv.imwrite('clahe_2.jpg', cl1)
plt.figure(figsize=(6, 6))
plt.imshow(cl1, cmap='gray')
plt.axis('off') # Hide axis plt.show()
print('Image written to clahe_2.jpg')

Image written to clahe_2.jpg

```



### #Practical-6-Program for Edge Detection

```

import cv2
import matplotlib.pyplot as plt
img = cv2.imread('cr7.jpg')
img_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
gray = cv2.cvtColor(img_rgb, cv2.COLOR_RGB2GRAY)
ret, thresh = cv2.threshold(gray, 125, 255, 0)
contours, hierarchy = cv2.findContours(thresh, cv2.RETR_TREE,
cv2.CHAIN_APPROX_SIMPLE)
copy_img = img_rgb.copy()
cv2.drawContours(copy_img, contours, -1, (0, 0, 255), 2)
fig, axes = plt.subplots(1, 3, figsize=(20, 10))
axes[0].imshow(img_rgb)
axes[0].set_title('Original Image')
axes[0].axis('off')
axes[1].imshow(gray, cmap='gray')
axes[1].set_title('Grayscale Image')
axes[1].axis('off')
axes[2].imshow(copy_img)
axes[2].set_title('Image with Contours')

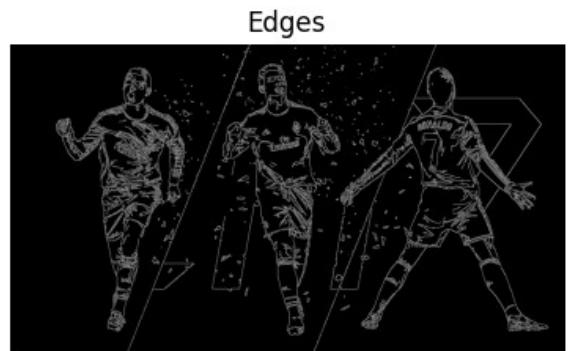
```

```
axes[2].axis('off')
plt.show()
```



### #Practical-7-Program for Line Detection

```
import cv2
import matplotlib.pyplot as plt
img = cv2.imread('cr7.jpg')
edges = cv2.Canny(img, 100, 200, apertureSize=3, L2gradient=True)
plt.imsave('test.png', edges, cmap='gray', format='png')
plt.figure(figsize=(10, 5))
plt.subplot(1, 2, 1)
plt.title('Original Image')
plt.imshow(cv2.cvtColor(img, cv2.COLOR_BGR2RGB))
plt.axis('off')
plt.subplot(1, 2, 2)
plt.title('Edges')
plt.imshow(edges, cmap='gray')
plt.axis('off')
plt.show()
print("Edges image saved as 'test.png'")
```



Edges image saved as 'test.png'

### #Practical-8-Programs using Scale Invariant Feature Transform (SIFT)

```
import cv2
import matplotlib.pyplot as plt
%matplotlib inline
```

```

img1 = cv2.imread('cr7.jpg')
gray1 = cv2.cvtColor(img1, cv2.COLOR_BGR2GRAY)
sift = cv2.SIFT_create()
keypoints_1, descriptors_1 = sift.detectAndCompute(gray1, None)
img_1 = cv2.drawKeypoints(img1, keypoints_1, None,
flags=cv2.DRAW_MATCHES_FLAGS_DRAW_RICH_KEYPOINTS)
plt.figure(figsize=(10, 10))
plt.imshow(cv2.cvtColor(img_1, cv2.COLOR_BGR2RGB))
plt.title('Image with SIFT Keypoints')
plt.axis('off')
plt.show()

```

Image with SIFT Keypoints

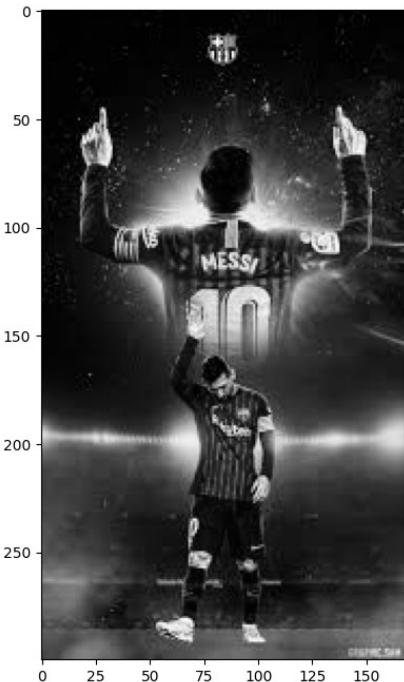
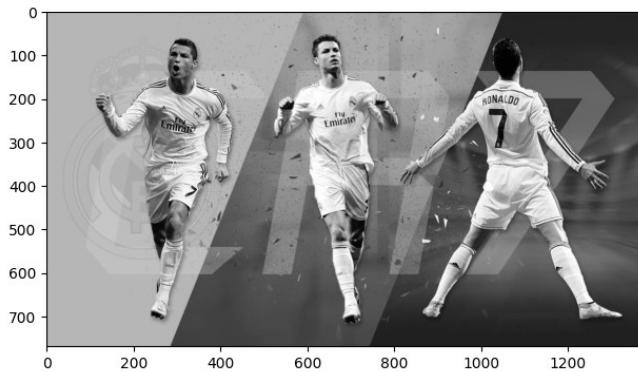


```

import cv2
import matplotlib.pyplot as plt
%matplotlib inline
img1 = cv2.imread('cr7.jpg')
img2 = cv2.imread('lm10.jpg')
img1 = cv2.cvtColor(img1, cv2.COLOR_BGR2GRAY)
img2 = cv2.cvtColor(img2, cv2.COLOR_BGR2GRAY)
figure,ax = plt.subplots(1,2,figsize=(16,8))
ax[0].imshow(img1,cmap='gray')
ax[1].imshow(img2,cmap='gray')

<matplotlib.image.AxesImage at 0x7e1dc04da980>

```



```
import cv2
import matplotlib.pyplot as plt
img1 = cv2.imread('cr7.jpg')
img2 = cv2.imread('lm10.jpg')
gray1 = cv2.cvtColor(img1, cv2.COLOR_BGR2GRAY)
gray2 = cv2.cvtColor(img2, cv2.COLOR_BGR2GRAY)
orb = cv2.ORB_create()
keypoints_1, descriptors_1 = orb.detectAndCompute(gray1, None)
keypoints_2, descriptors_2 = orb.detectAndCompute(gray2, None)
bf = cv2.BFMatcher(cv2.NORM_HAMMING, crossCheck=True)
matches = bf.match(descriptors_1, descriptors_2)
matches = sorted(matches, key=lambda x: x.distance)
img_matches = cv2.drawMatches(img1, keypoints_1, img2, keypoints_2,
matches[:50], None, flags=cv2.DrawMatchesFlags_NOT_DRAW_SINGLE_POINTS)
plt.figure(figsize=(12, 6))
plt.imshow(cv2.cvtColor(img_matches, cv2.COLOR_BGR2RGB))
plt.title('ORB Feature Matching')
plt.axis('off')
plt.show()
```

### ORB Feature Matching



```
import cv2
import matplotlib.pyplot as plt
img1 = cv2.imread('cr7.jpg')
img2 = cv2.imread('lm10.jpg')
gray1 = cv2.cvtColor(img1, cv2.COLOR_BGR2GRAY)
gray2 = cv2.cvtColor(img2, cv2.COLOR_BGR2GRAY)
sift = cv2.SIFT_create()
keypoints_1, descriptors_1 = sift.detectAndCompute(gray1, None)
keypoints_2, descriptors_2 = sift.detectAndCompute(gray2, None)

bf = cv2.BFMatcher()
matches = bf.knnMatch(descriptors_1, descriptors_2, k=2)

good_matches = []
for m, n in matches:
    if m.distance < 0.75 * n.distance: good_matches.append(m)

img_matches = cv2.drawMatches(img1, keypoints_1, img2, keypoints_2,
good_matches, None, flags=cv2.DrawMatchesFlags_NOT_DRAW_SINGLE_POINTS)
fig, axes = plt.subplots(1, 3, figsize=(15, 5))
axes[0].imshow(cv2.cvtColor(img1, cv2.COLOR_BGR2RGB))
axes[0].set_title('Image 1')
axes[0].axis('off')
axes[1].imshow(cv2.cvtColor(img_matches, cv2.COLOR_BGR2RGB))
axes[1].set_title('Matches')
axes[1].axis('off')
axes[2].imshow(cv2.cvtColor(img2, cv2.COLOR_BGR2RGB))
axes[2].set_title('Image 2')
```

```
axes[2].axis('off')
plt.tight_layout()
plt.show()
```

Image 1



Matches

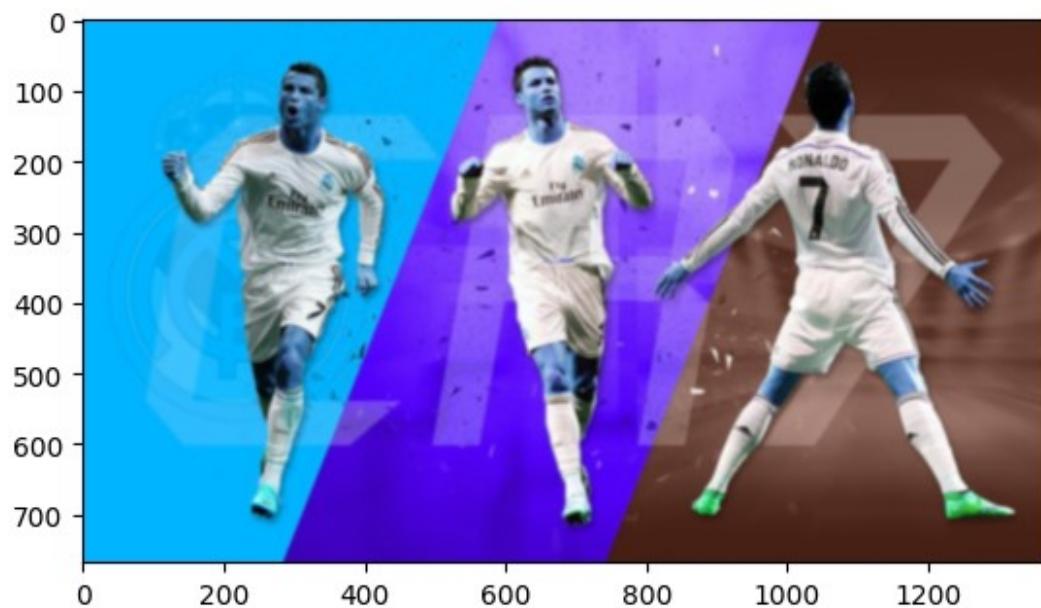


Image 2



```
def gaussian_blur(image, kernel_size=(5, 5), sigma=0): return
cv2.GaussianBlur(image, kernel_size, sigma)
image = cv2.imread('cr7.jpg')
blurred_image = gaussian_blur(image, kernel_size=(11, 11), sigma=0)
plt.imshow(blurred_image)

<matplotlib.image.AxesImage at 0x7e1dc05bfb20>
```



```
from google.colab import files
import cv2
import numpy as np

# Upload video file
uploaded = files.upload()

# Assuming the uploaded file is named 'video.mp4'
cap = cv2.VideoCapture('Siiuu.mp4')

last_mean = 0

while cap.isOpened():
    ret, frame = cap.read()

    if not ret:
        print("Failed to grab frame or end of video")
        break

    # Display the resulting frame
    # In Colab, you would typically save and display the frame using
cv2_imshow
    from google.colab.patches import cv2_imshow
    cv2_imshow(frame)

    # Convert the frame to grayscale
    gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)

    # Calculate the mean difference
    result = np.abs(np.mean(gray) - last_mean)
    print(result)

    # Detect change
    if result > 0:
        print("Change detected")
    else:
        print("Change not detected")

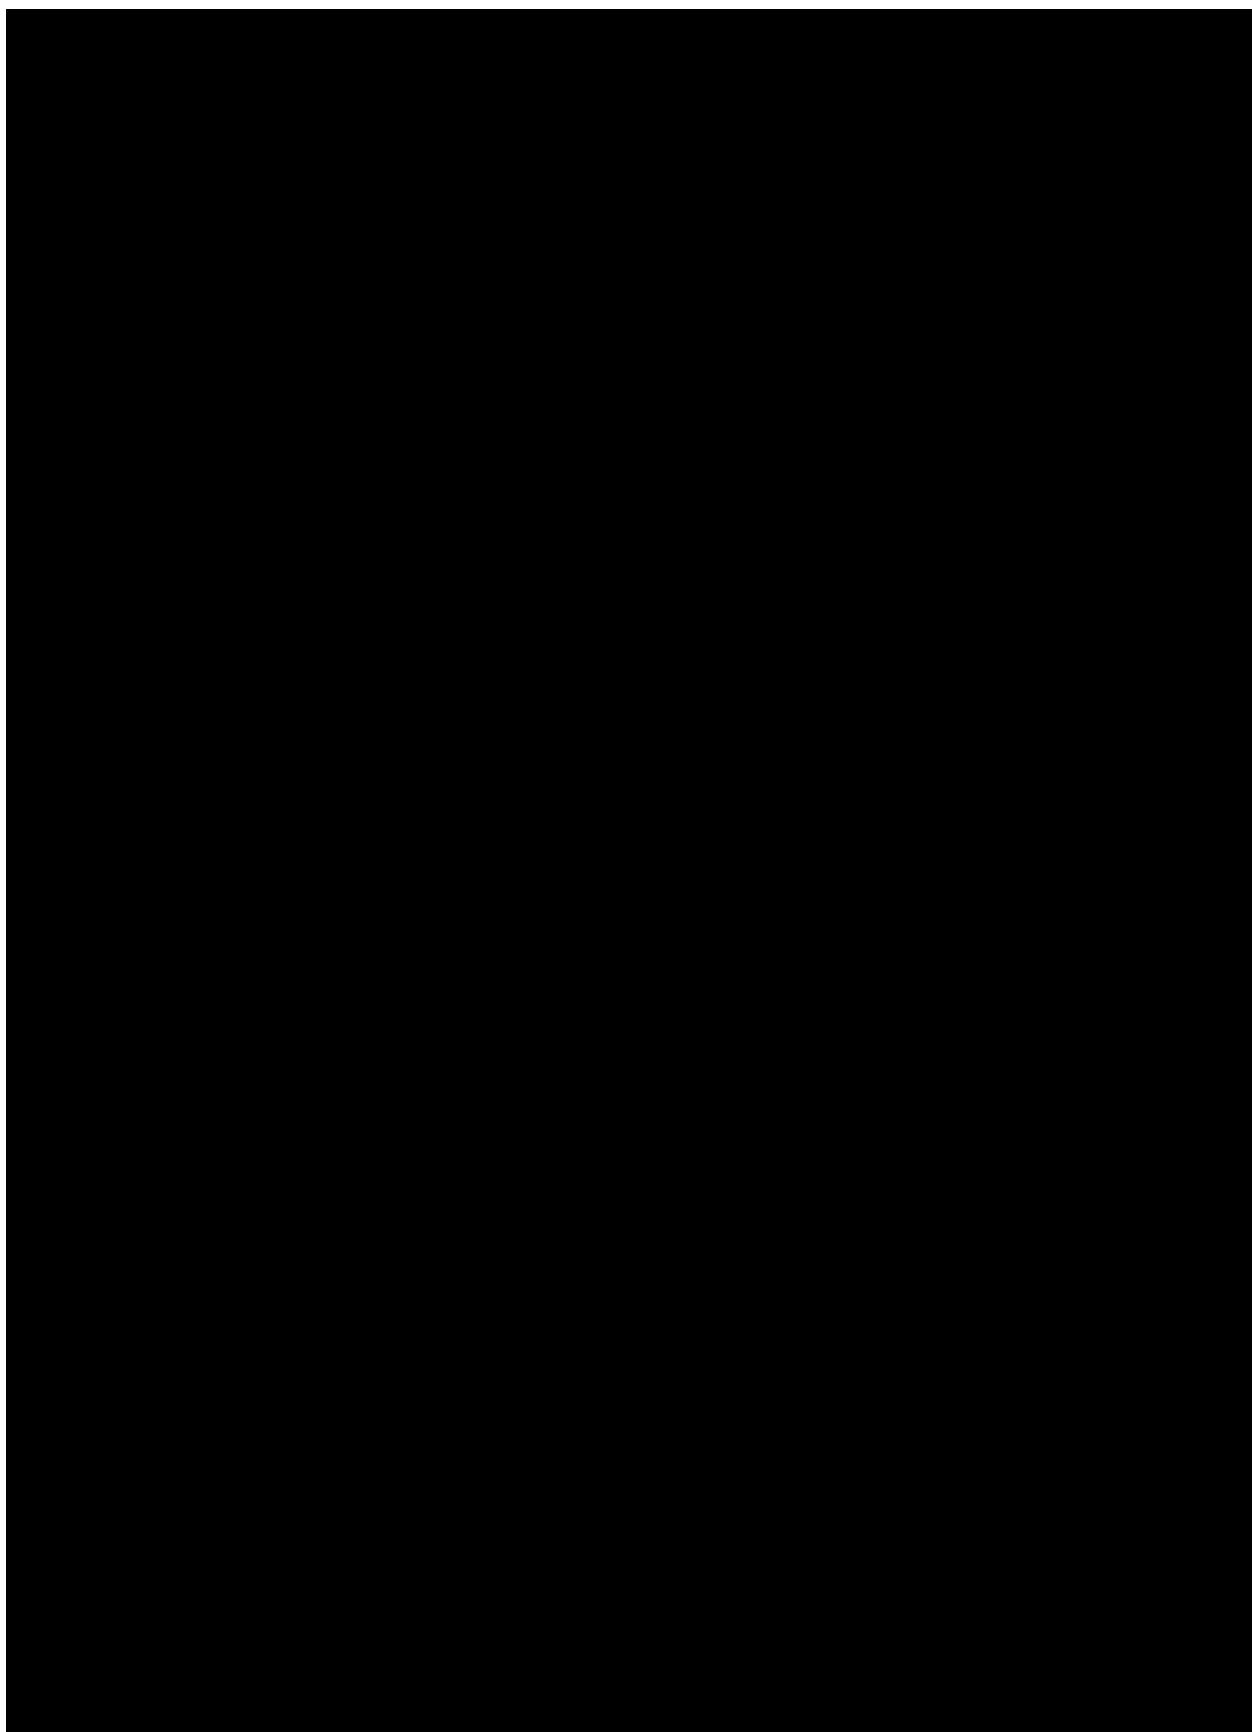
    # Update last_mean
    last_mean = np.mean(gray)

    # Wait for a while before showing the next frame
    if cv2.waitKey(1) & 0xFF == ord('q'):
        break

# When everything is done, release the capture
cap.release()
cv2.destroyAllWindows()

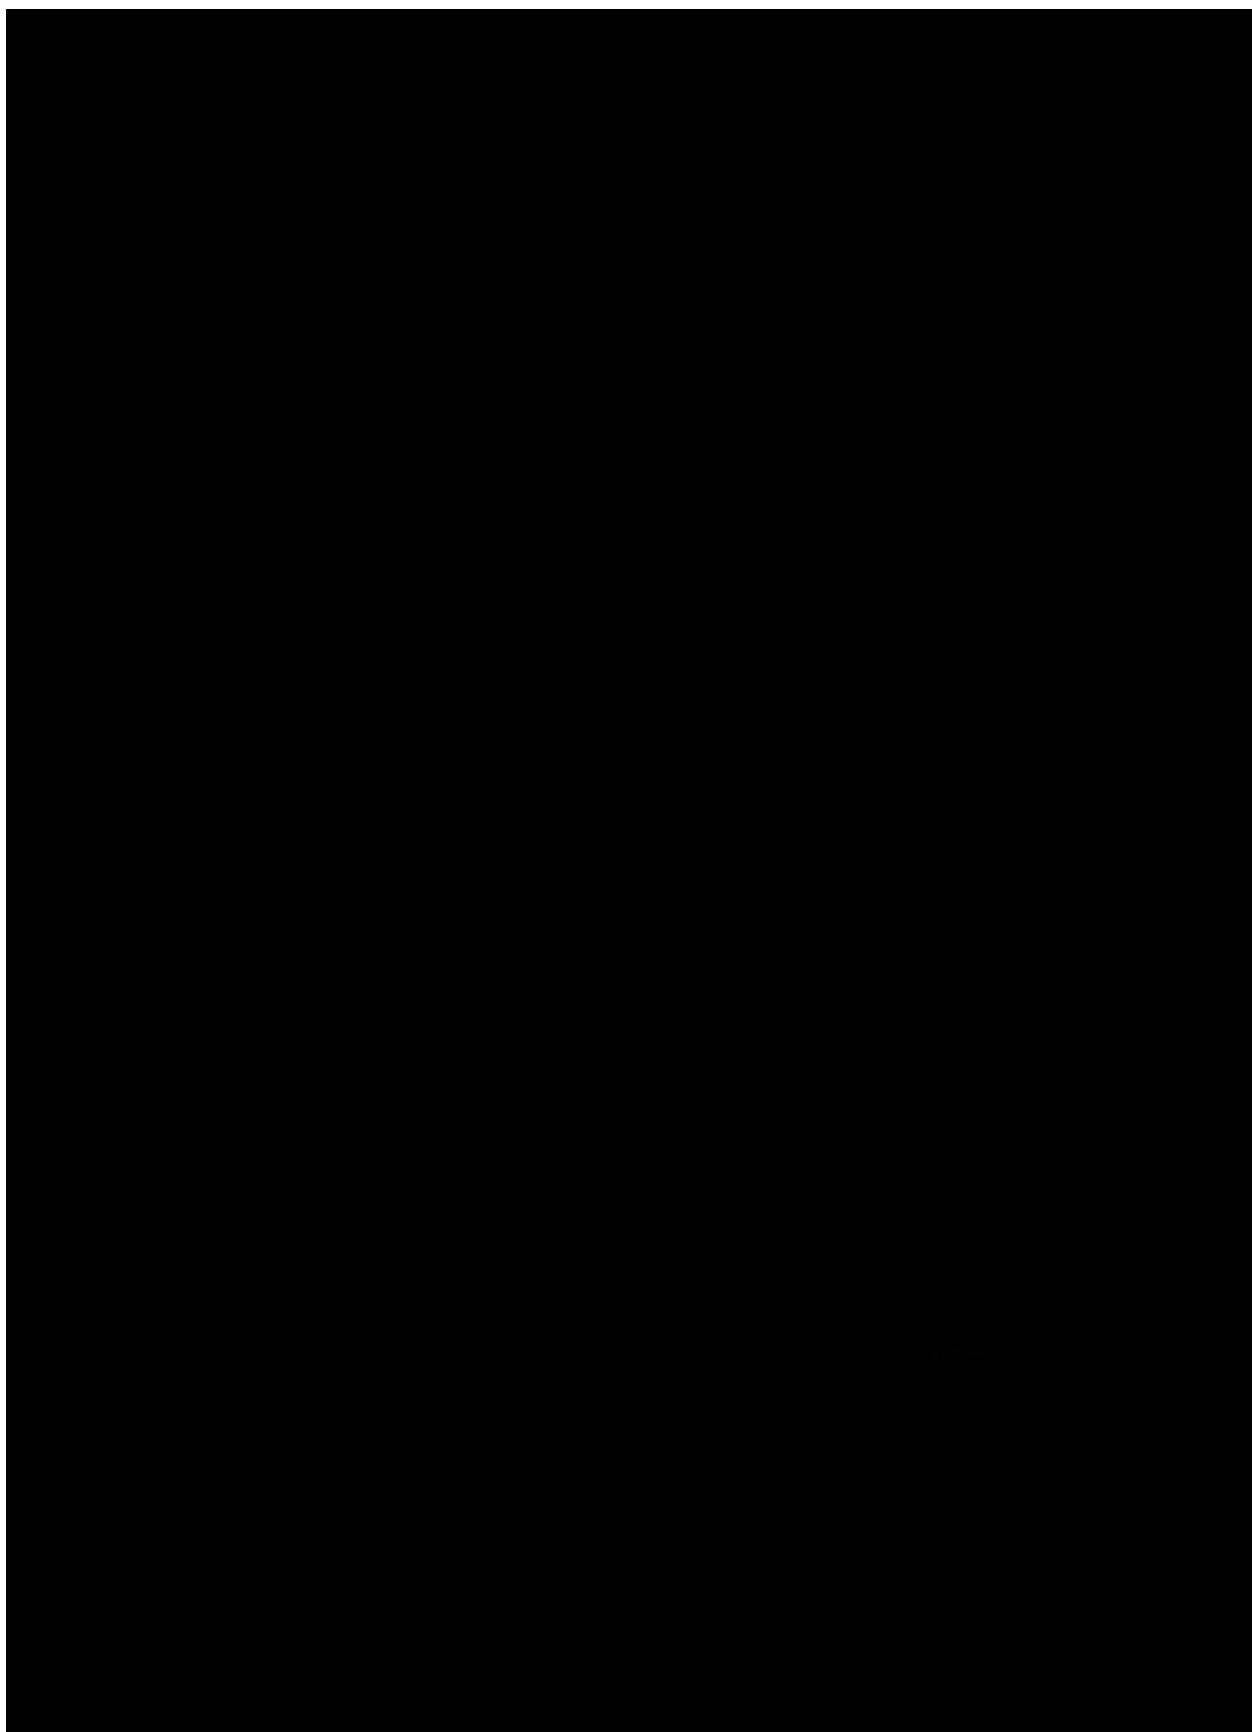
<IPython.core.display.HTML object>
```

Saving Siiuuu.mp4 to Siiuuu (1).mp4



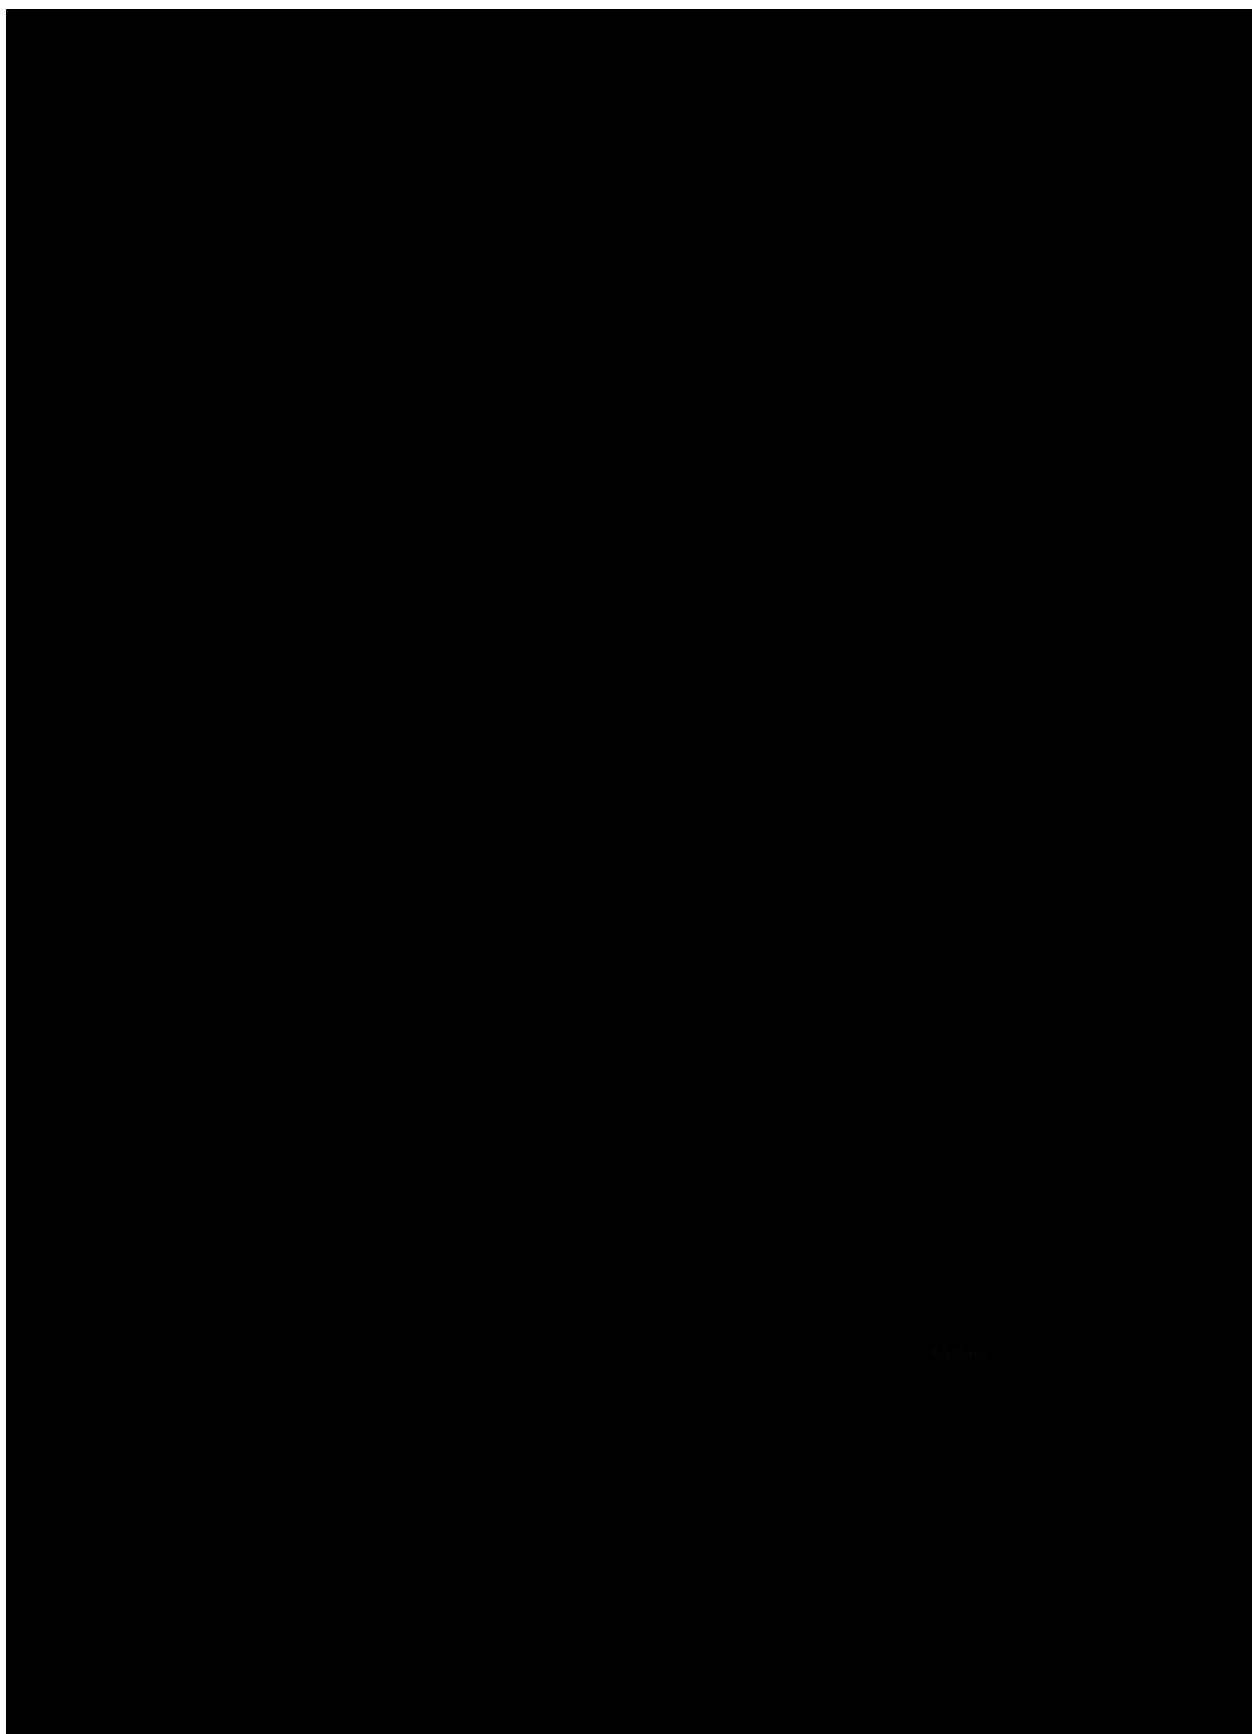
0.0

Change not detected



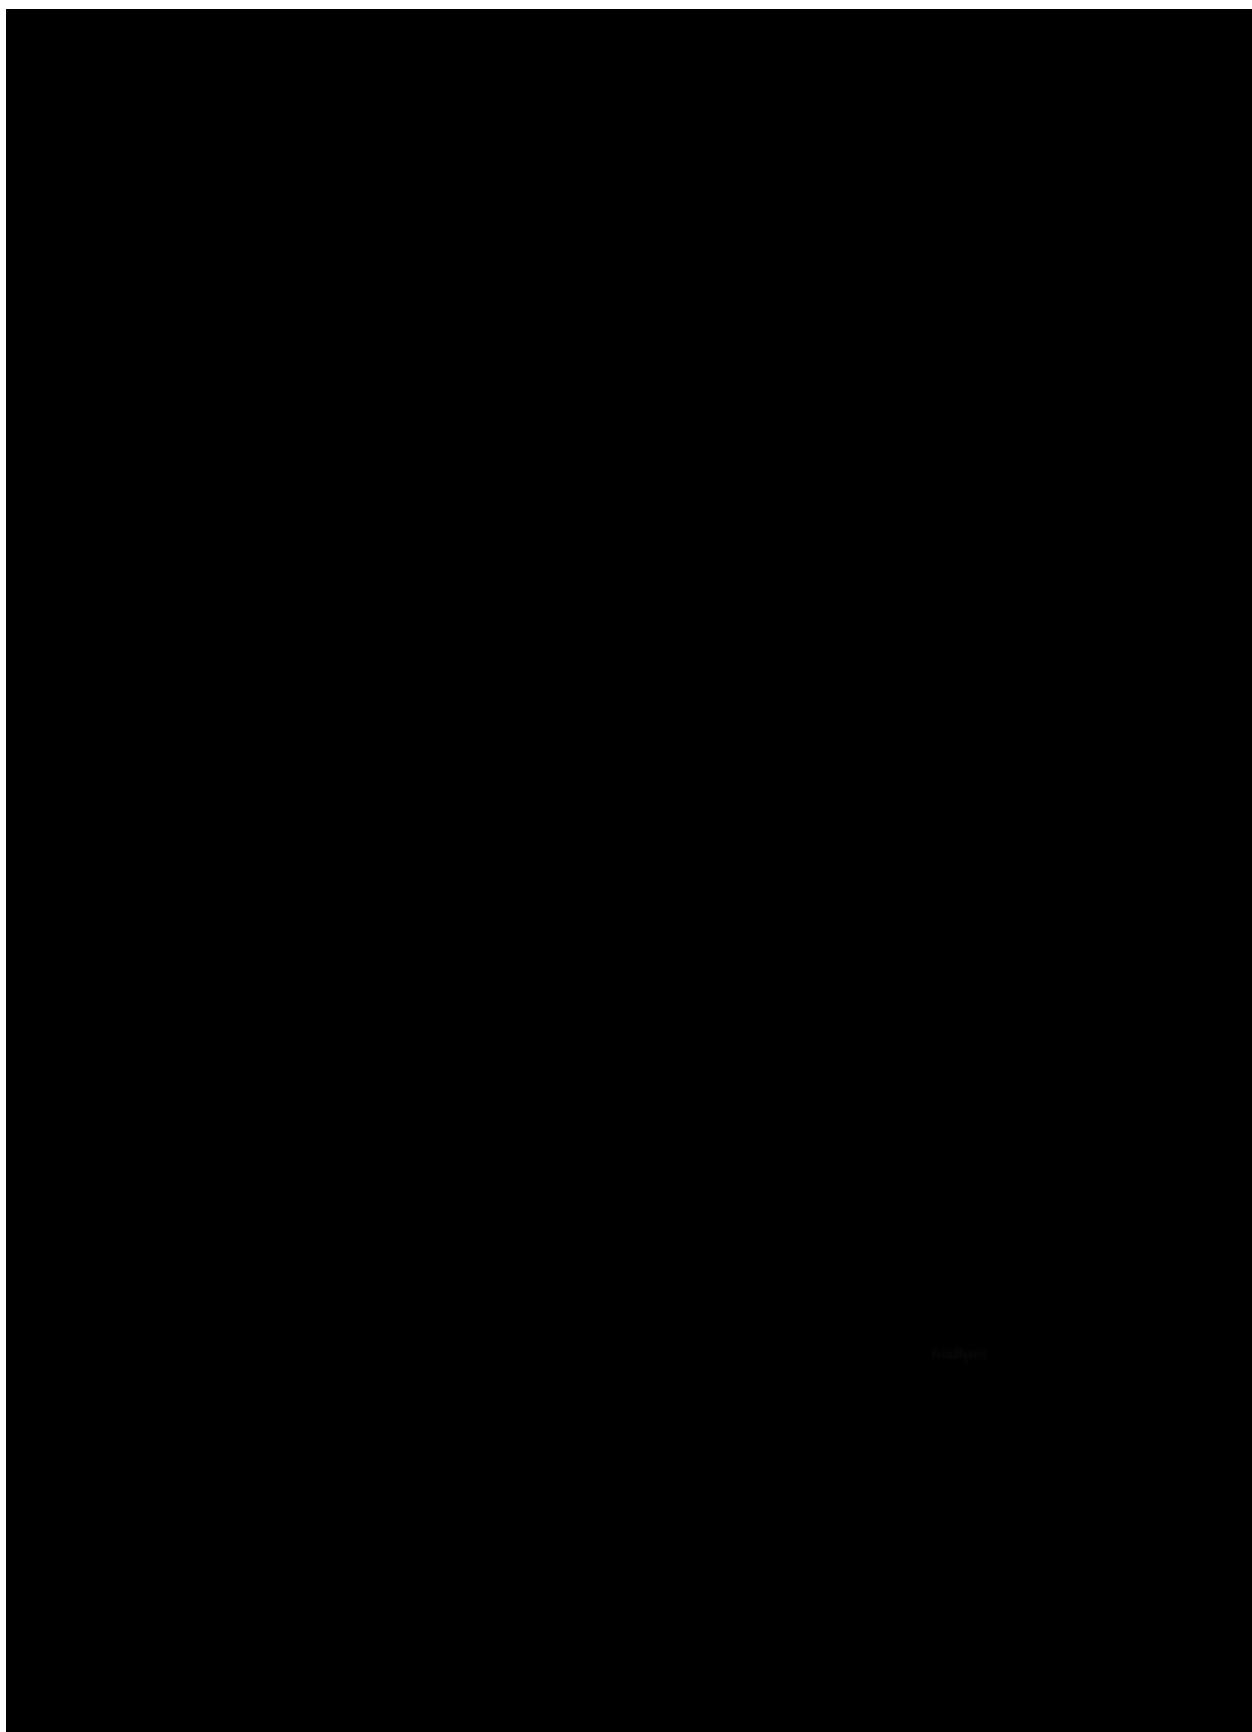
0.00038773148148148147

Change detected



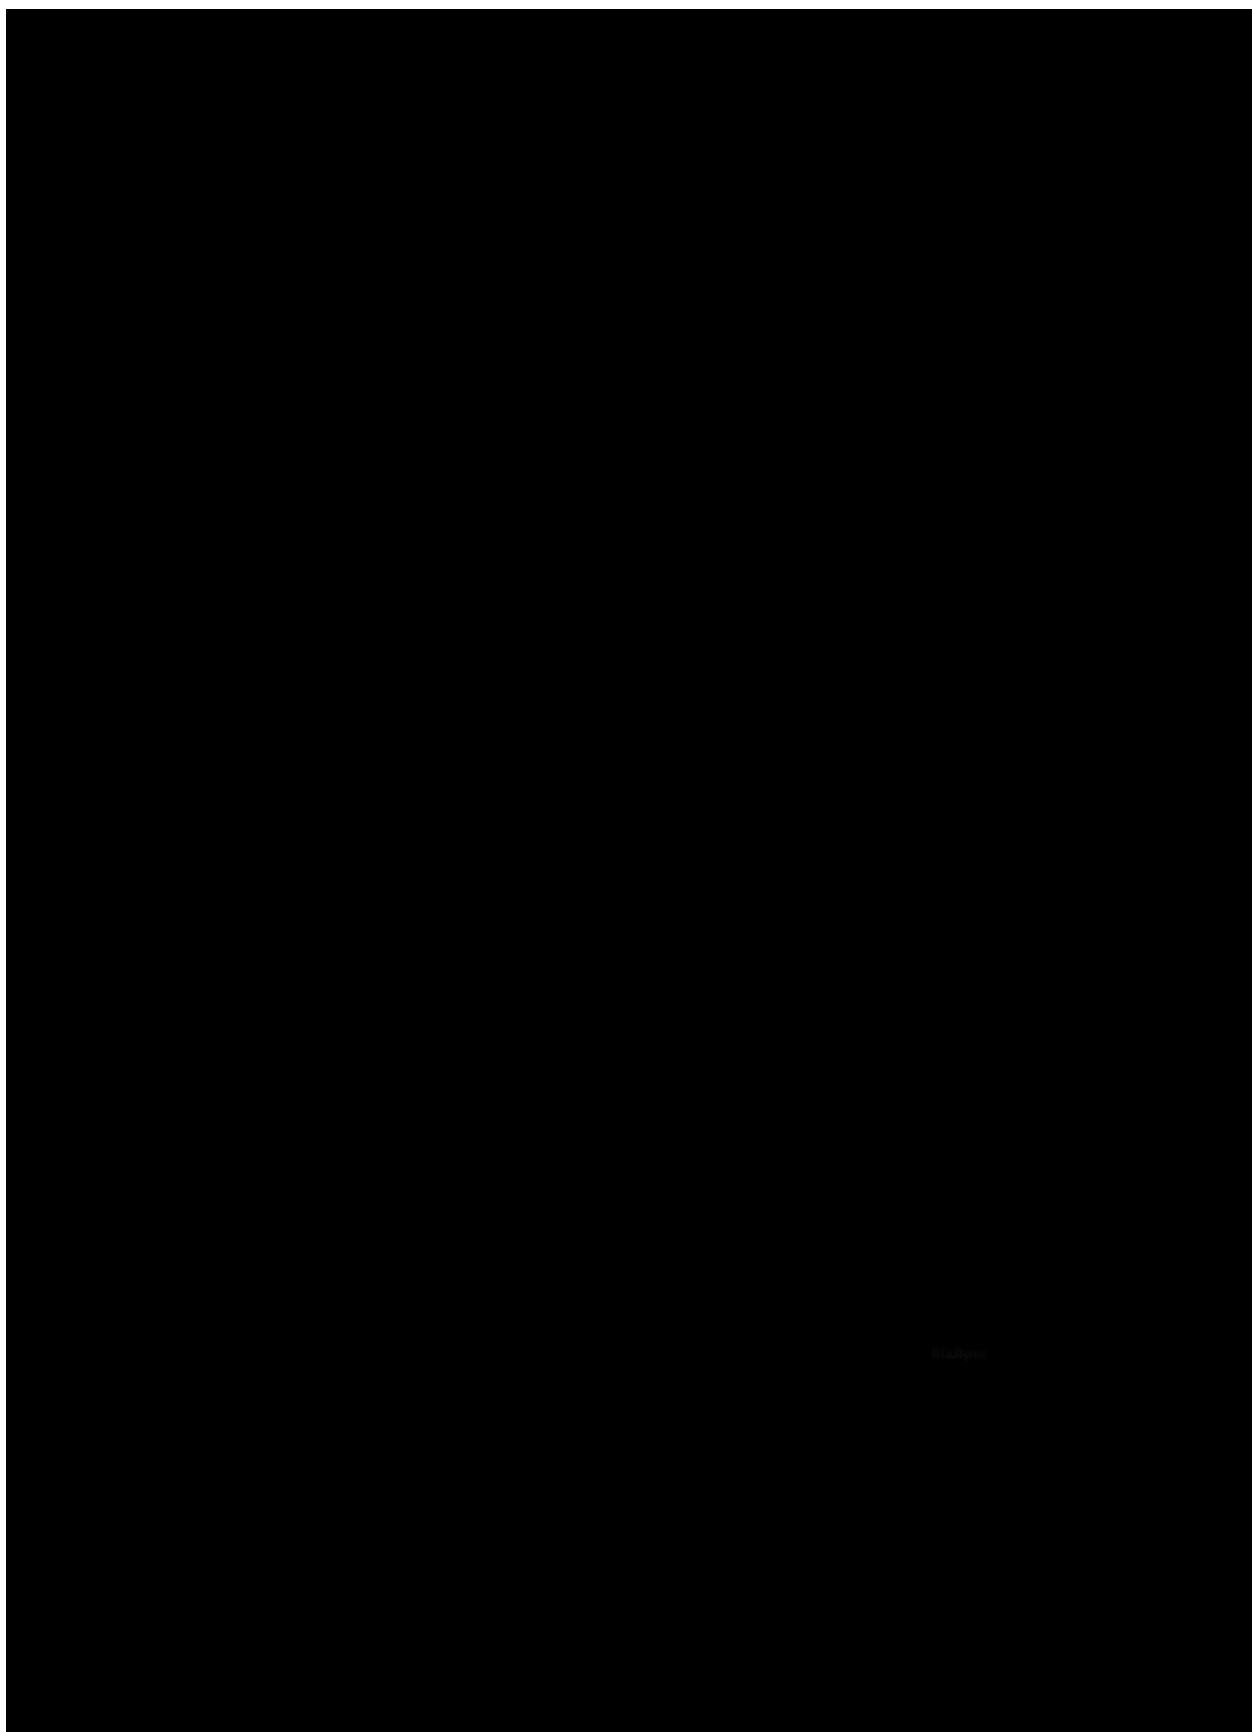
0.00024353780864197532

Change detected



0.0003014081790123457

Change detected



0.00022665895061728392

Change detected

*hizigoi*

0.0007325424382716049

Change detected

*hizligol*

0.0003149112654320988

Change detected

*hizligol*

0.0005135995370370372

Change detected

*hitzigoi*

0.0003583140432098764

Change detected

*hitzigoi*

0.0003993055555555555

Change detected

*hitzigoi*

0.00026620370370370383

Change detected

*hitzigoi*

0.00035927854938271574

Change detected

*hitzigol*

0.00018904320987654308

Change detected

*hitzigol*

0.0005671296296296301

Change detected

*hitzigol*

0.0003896604938271606

Change detected

*hitzigol*

0.00012827932098765423

Change detected

*hitzigol*

7.908950617283896e-05

Change detected

*hitzigol*

0.000590277777777781

Change detected

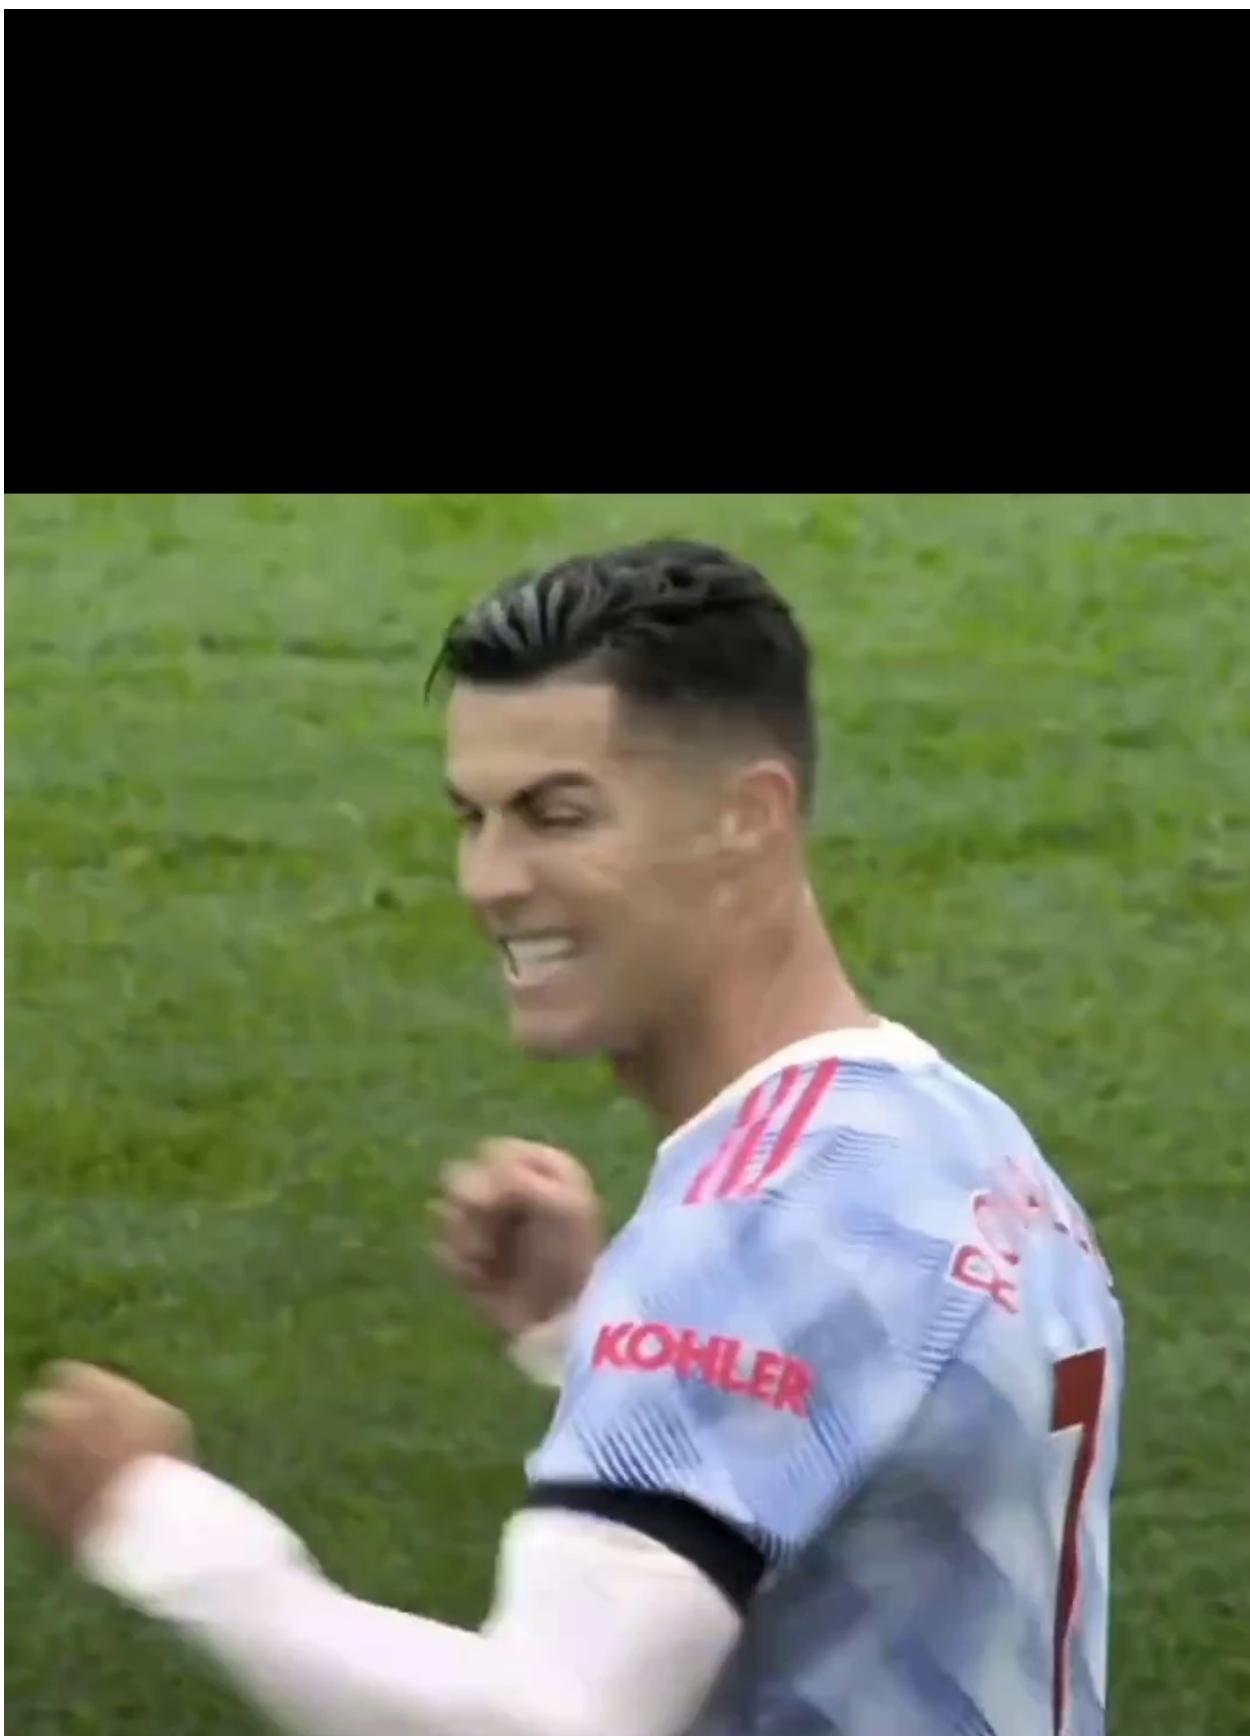
*hitzigol*

6.269290123456835e-05  
Change detected



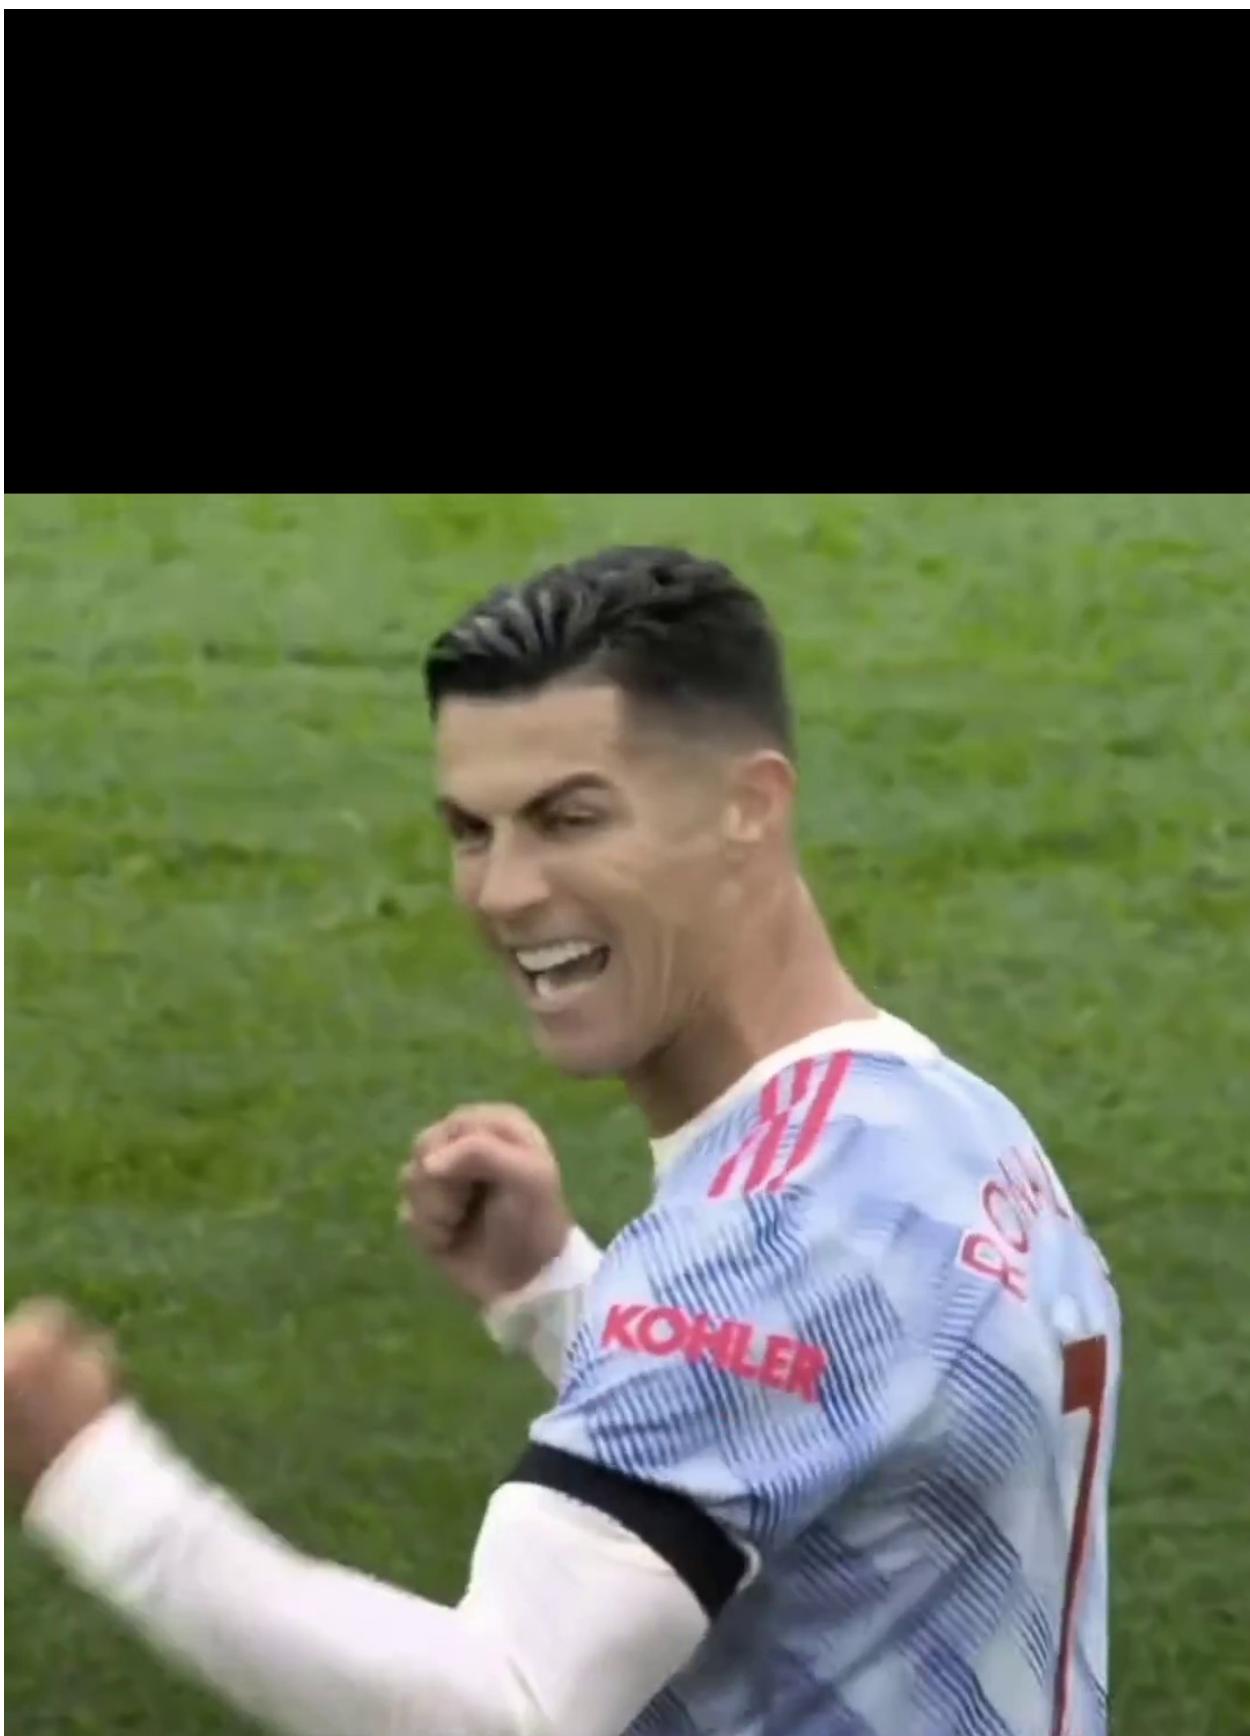
68.73231047453703

Change detected



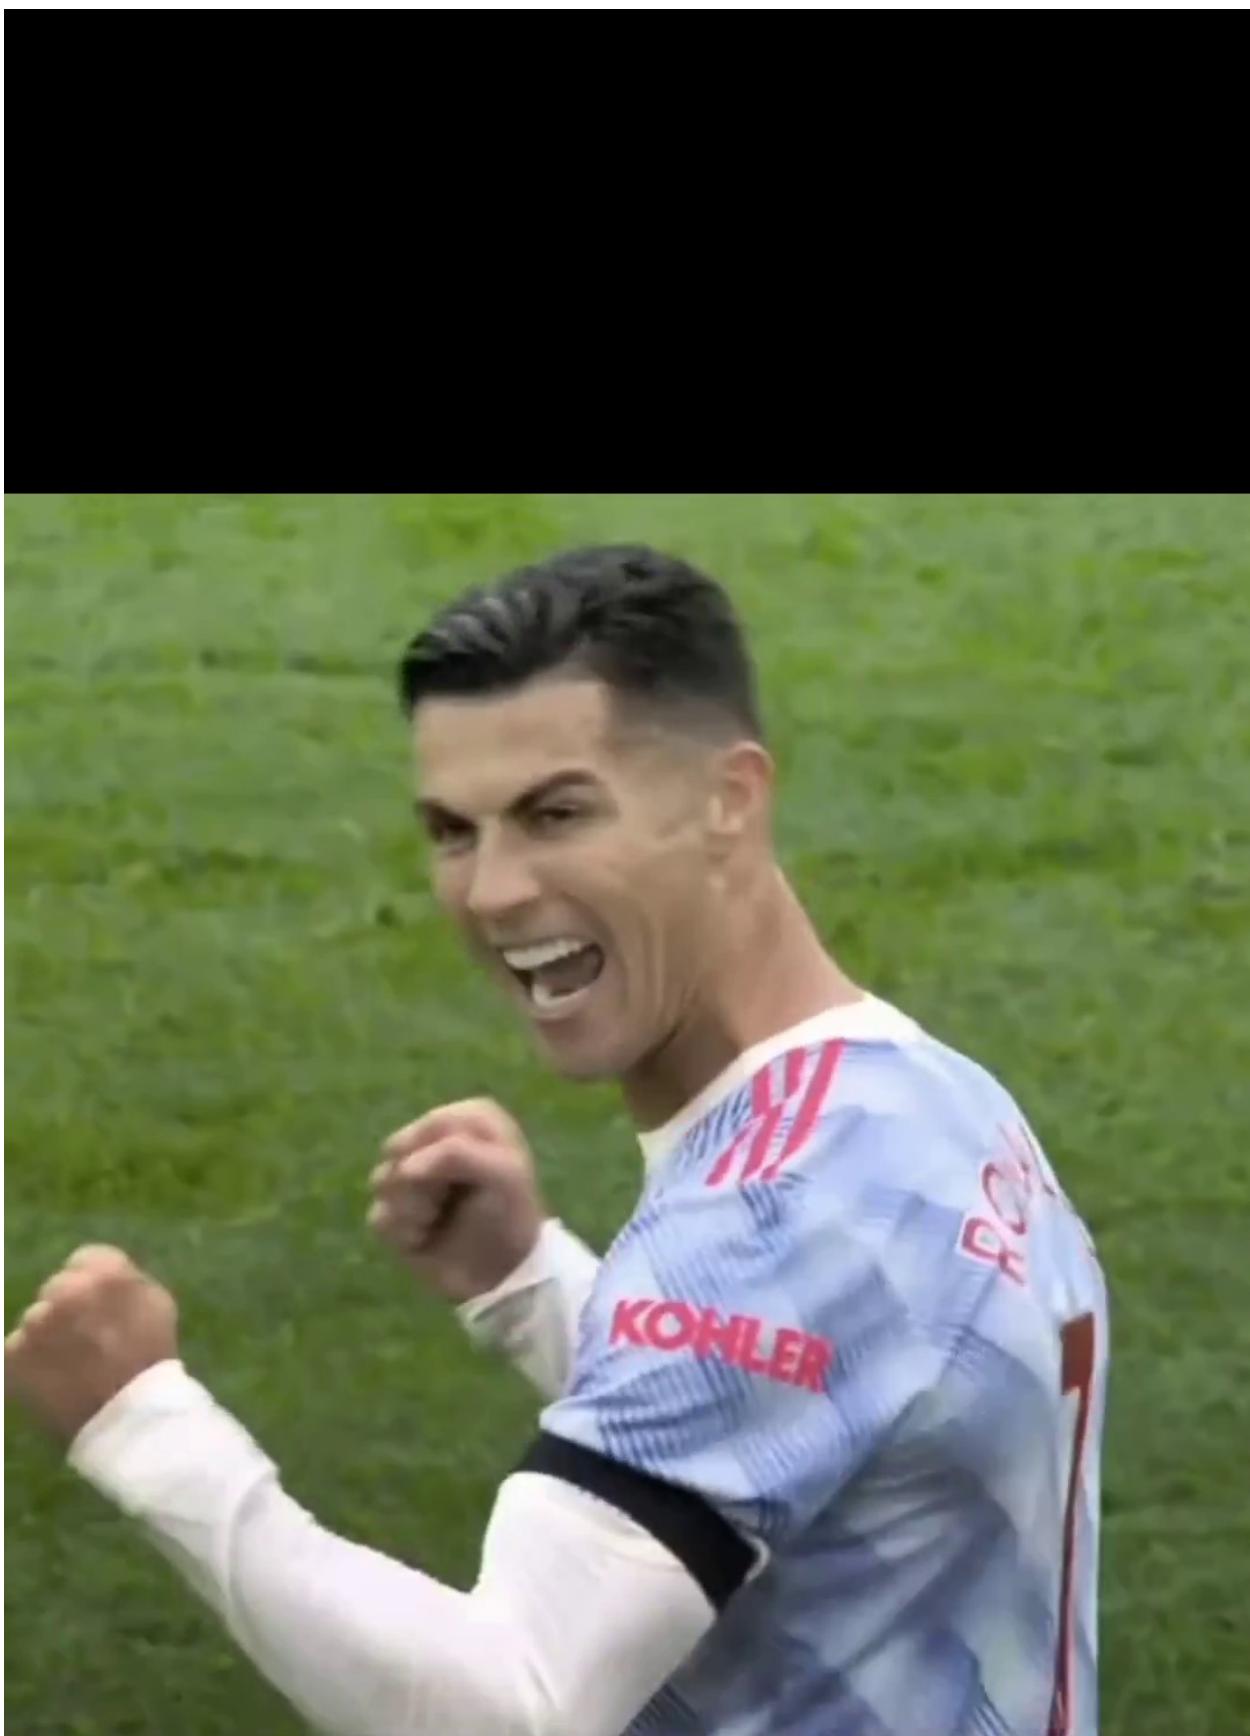
1.6368822337962996

Change detected



0.22717062114197972

Change detected



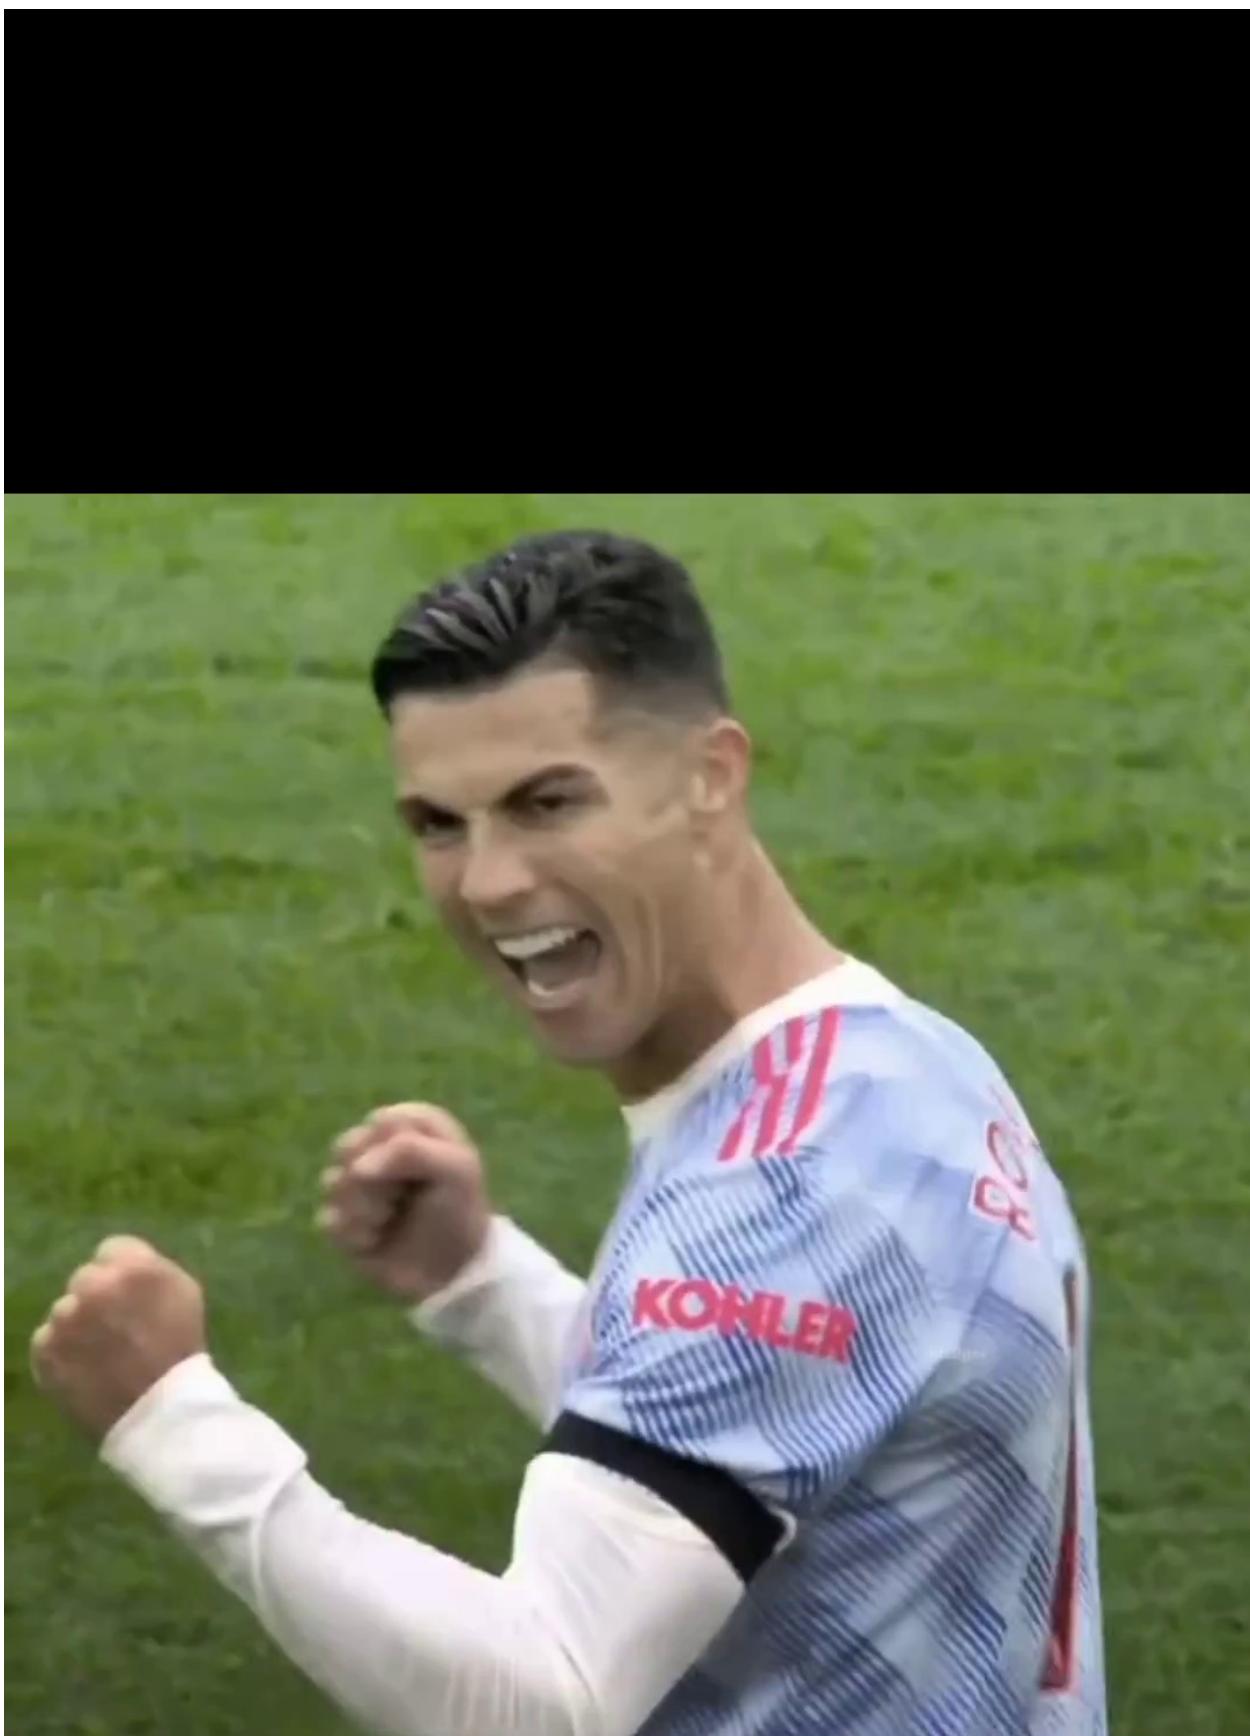
0.16228684413580652

Change detected



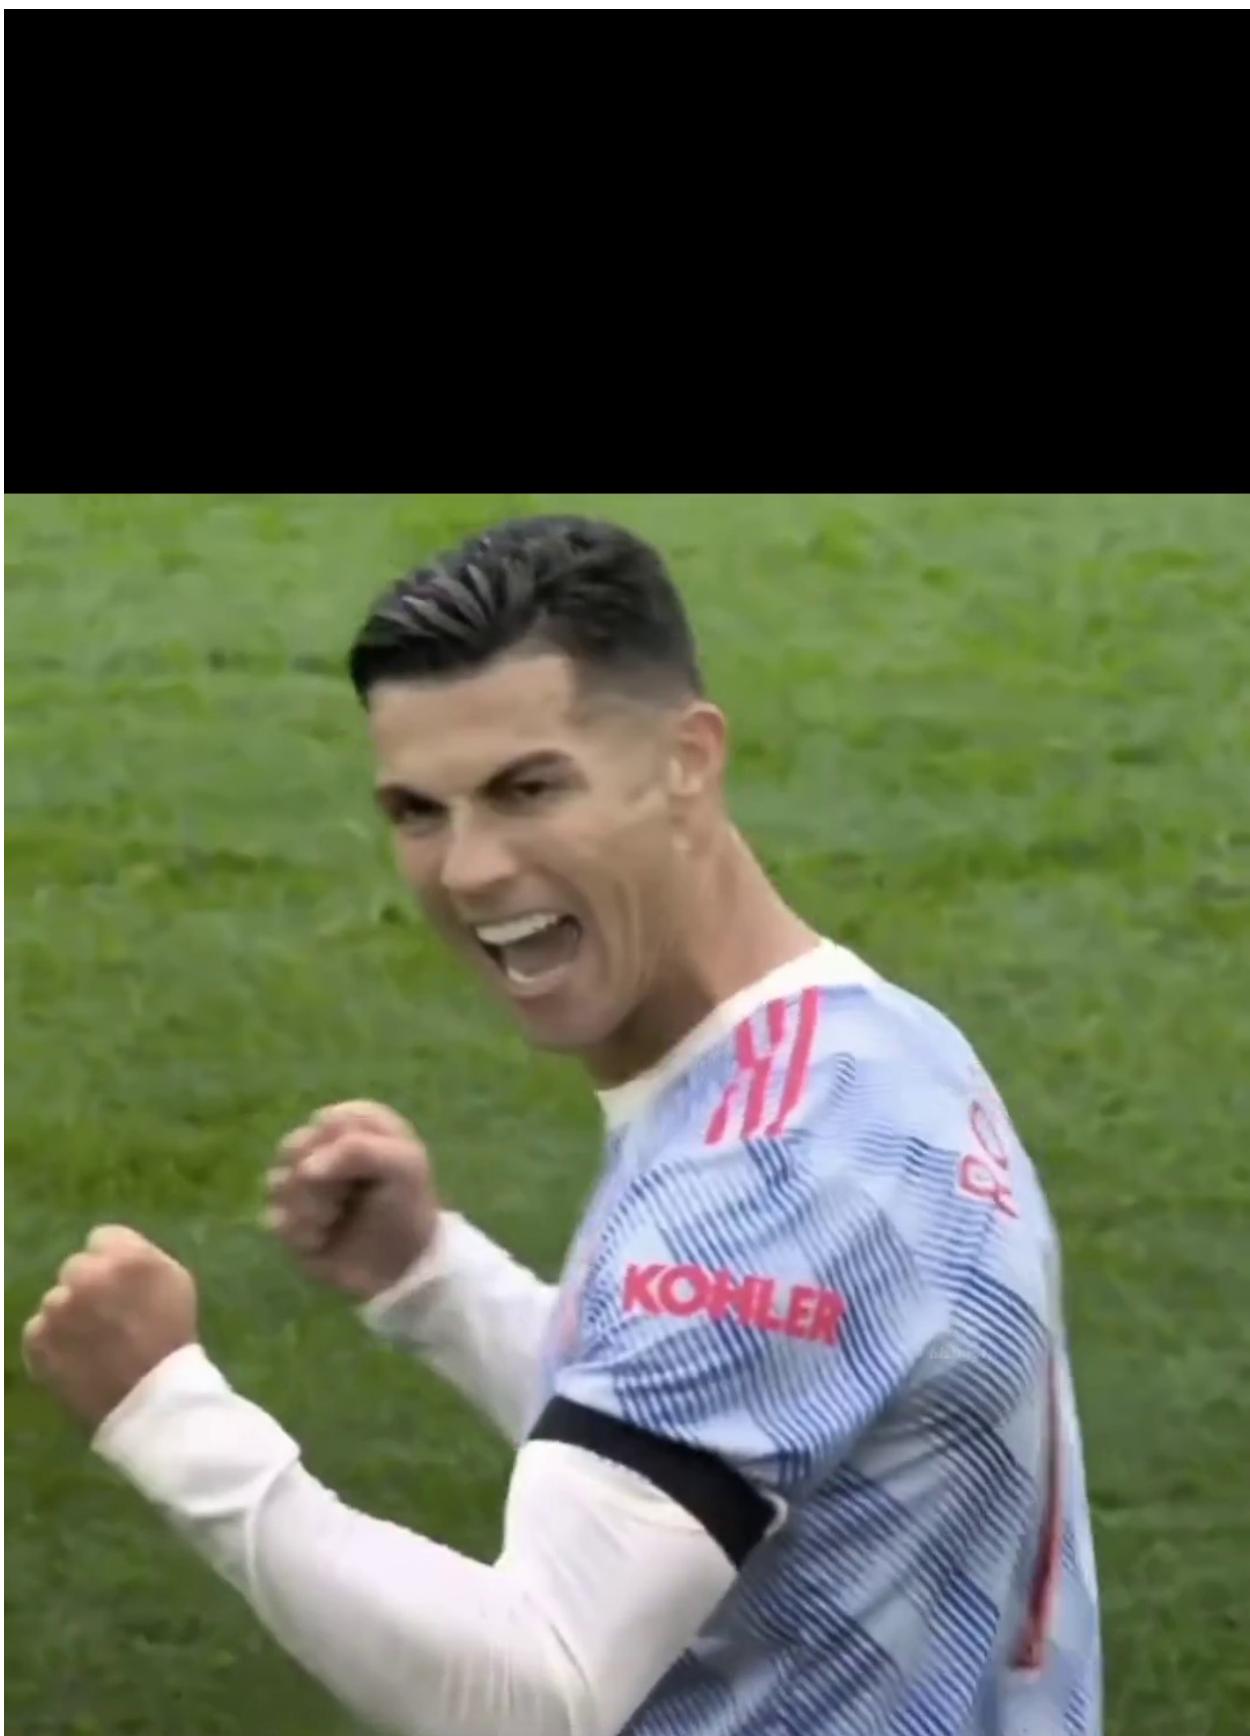
0.17955295138888516

Change detected



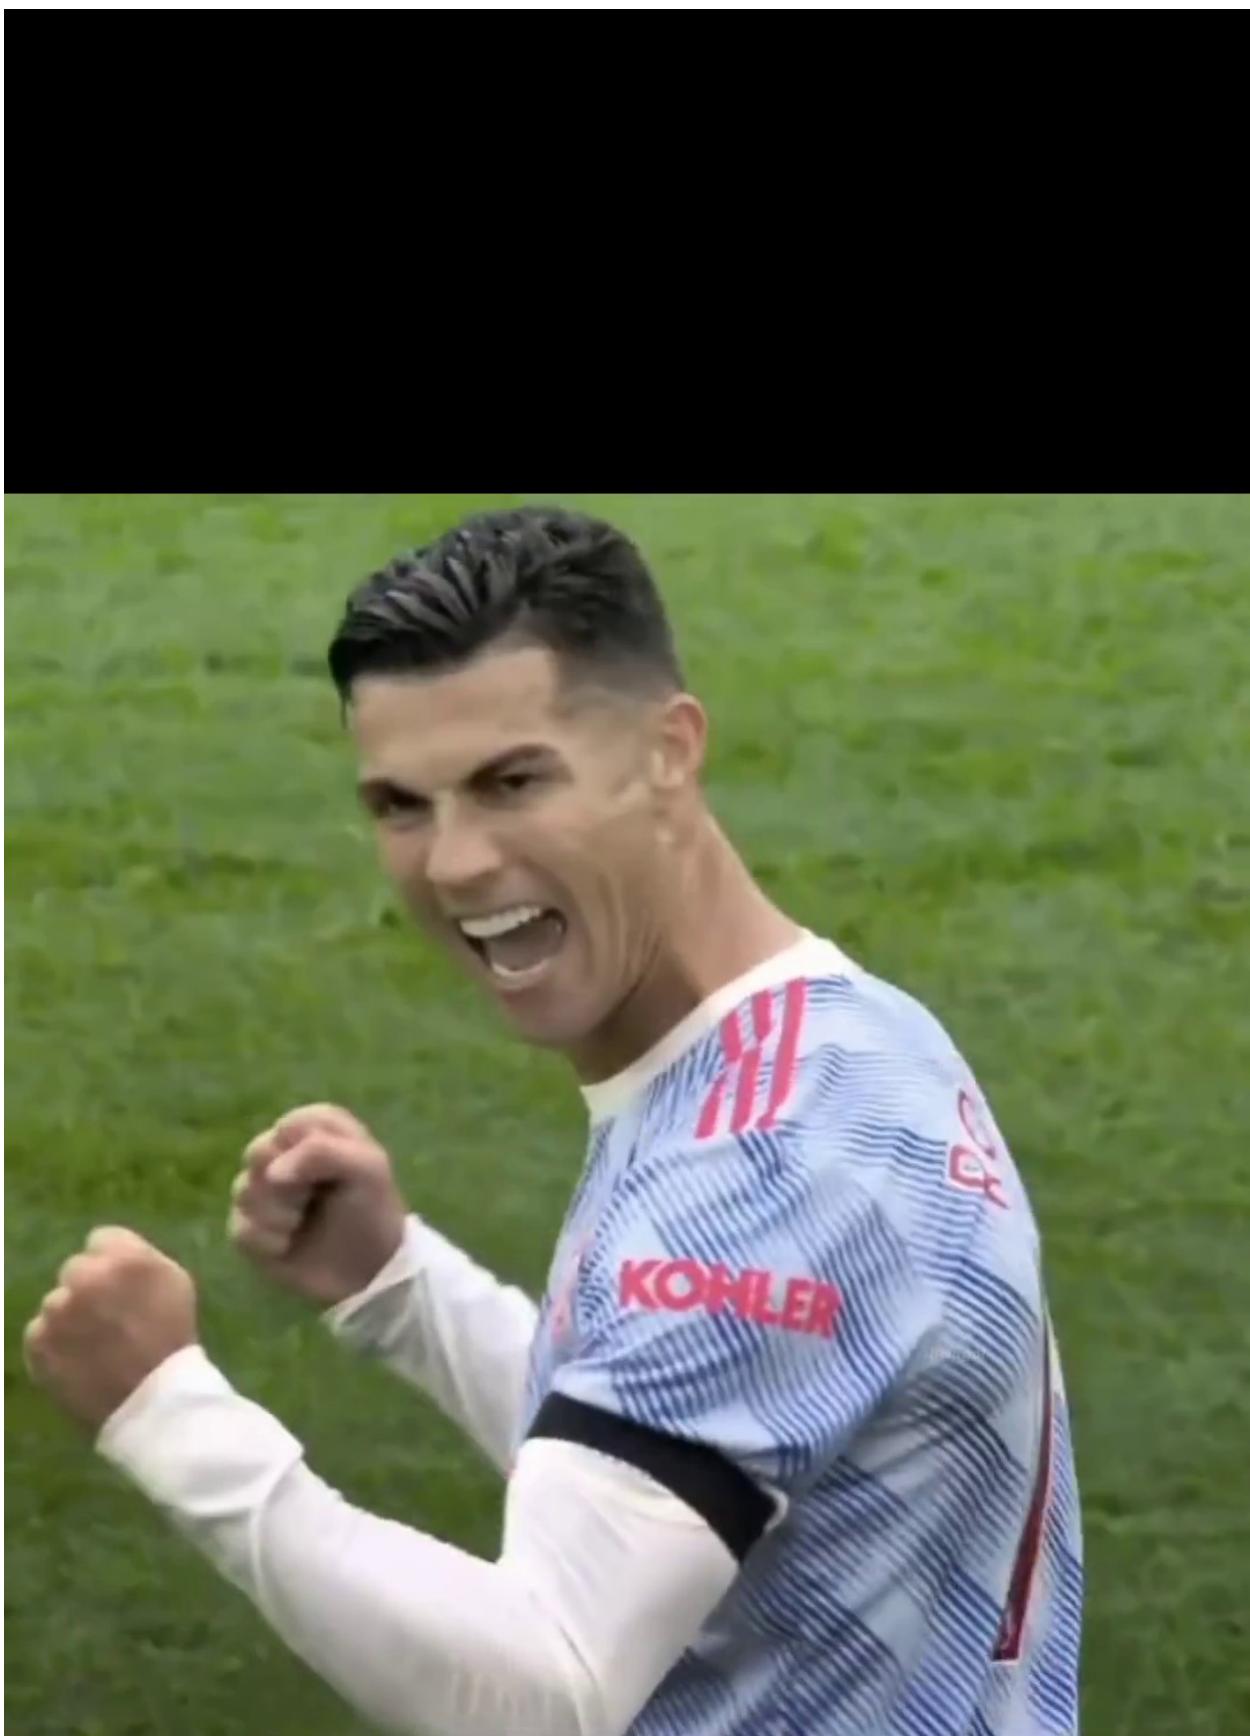
0.11328366126544154

Change detected



0.49167824074075384

Change detected



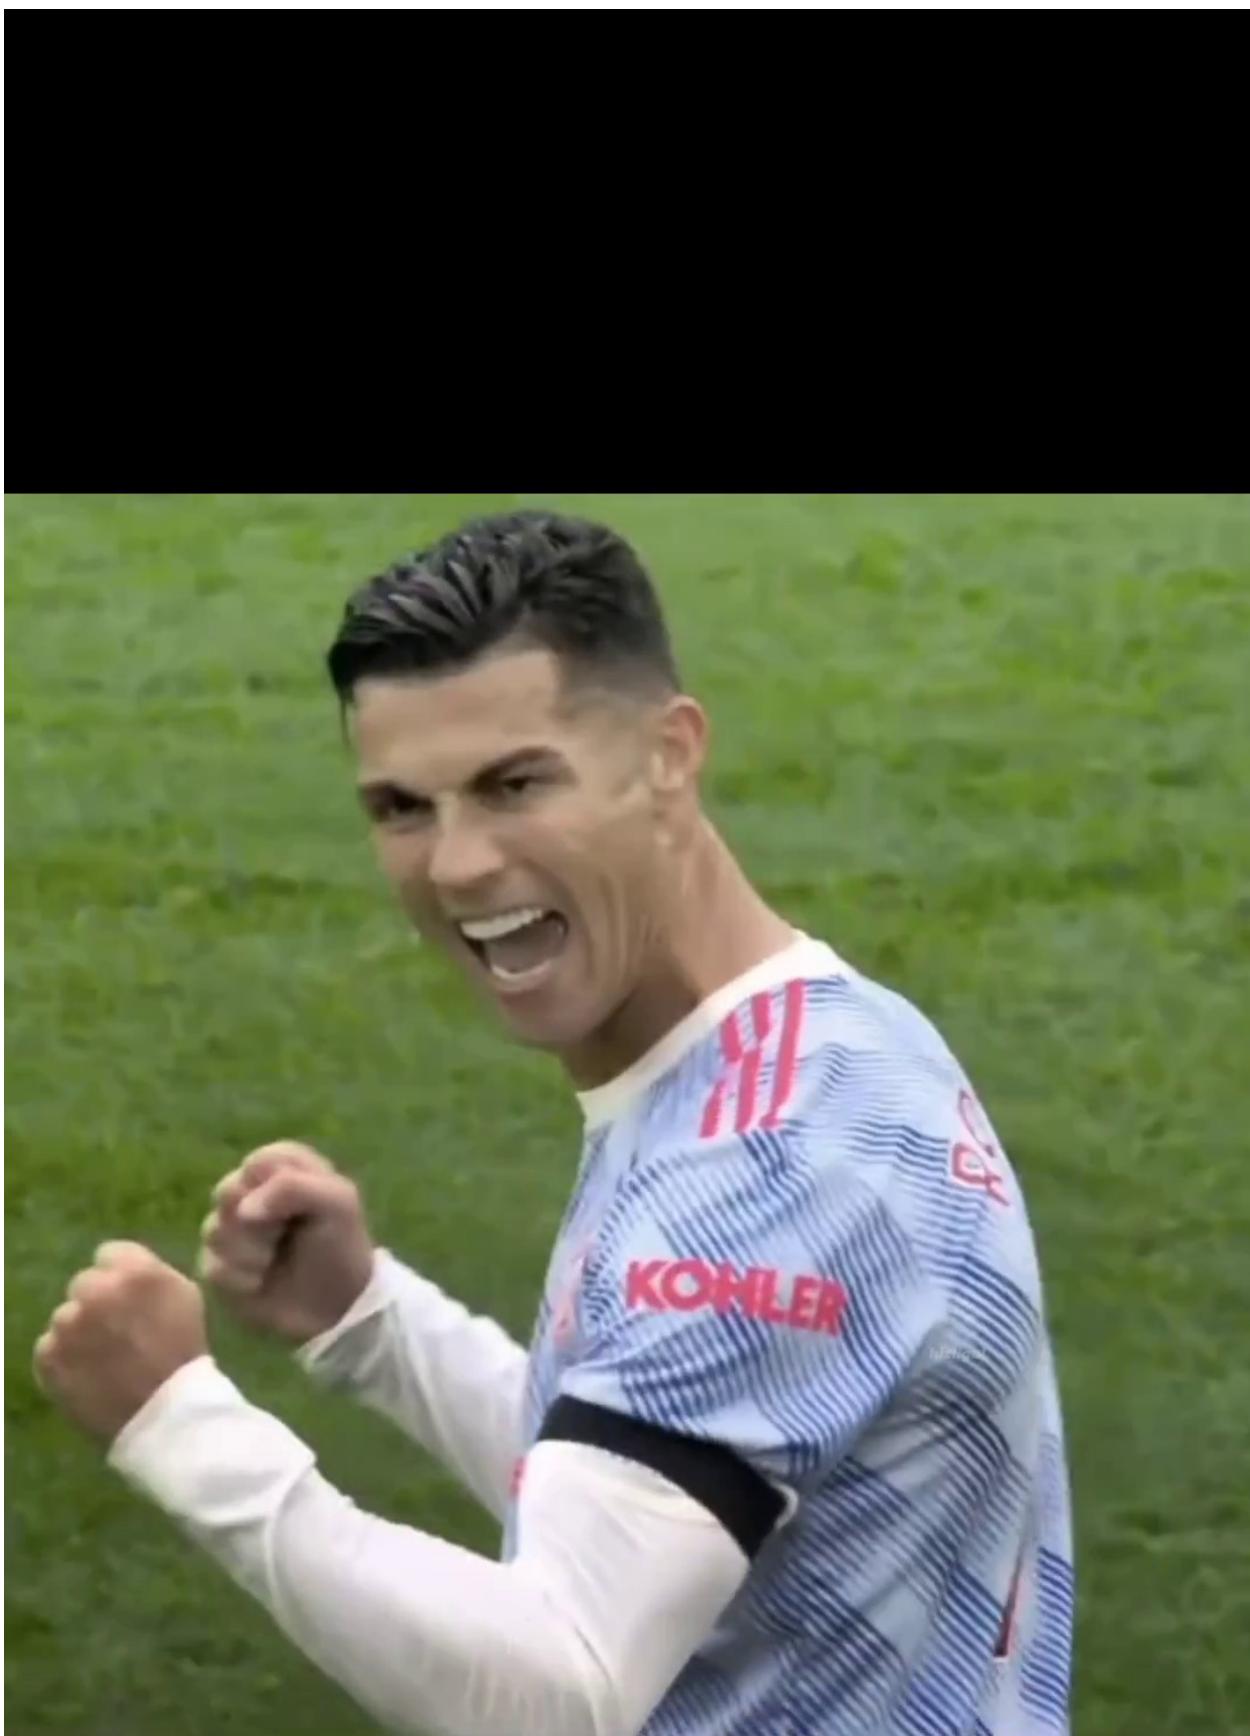
0.22325520833334167

Change detected



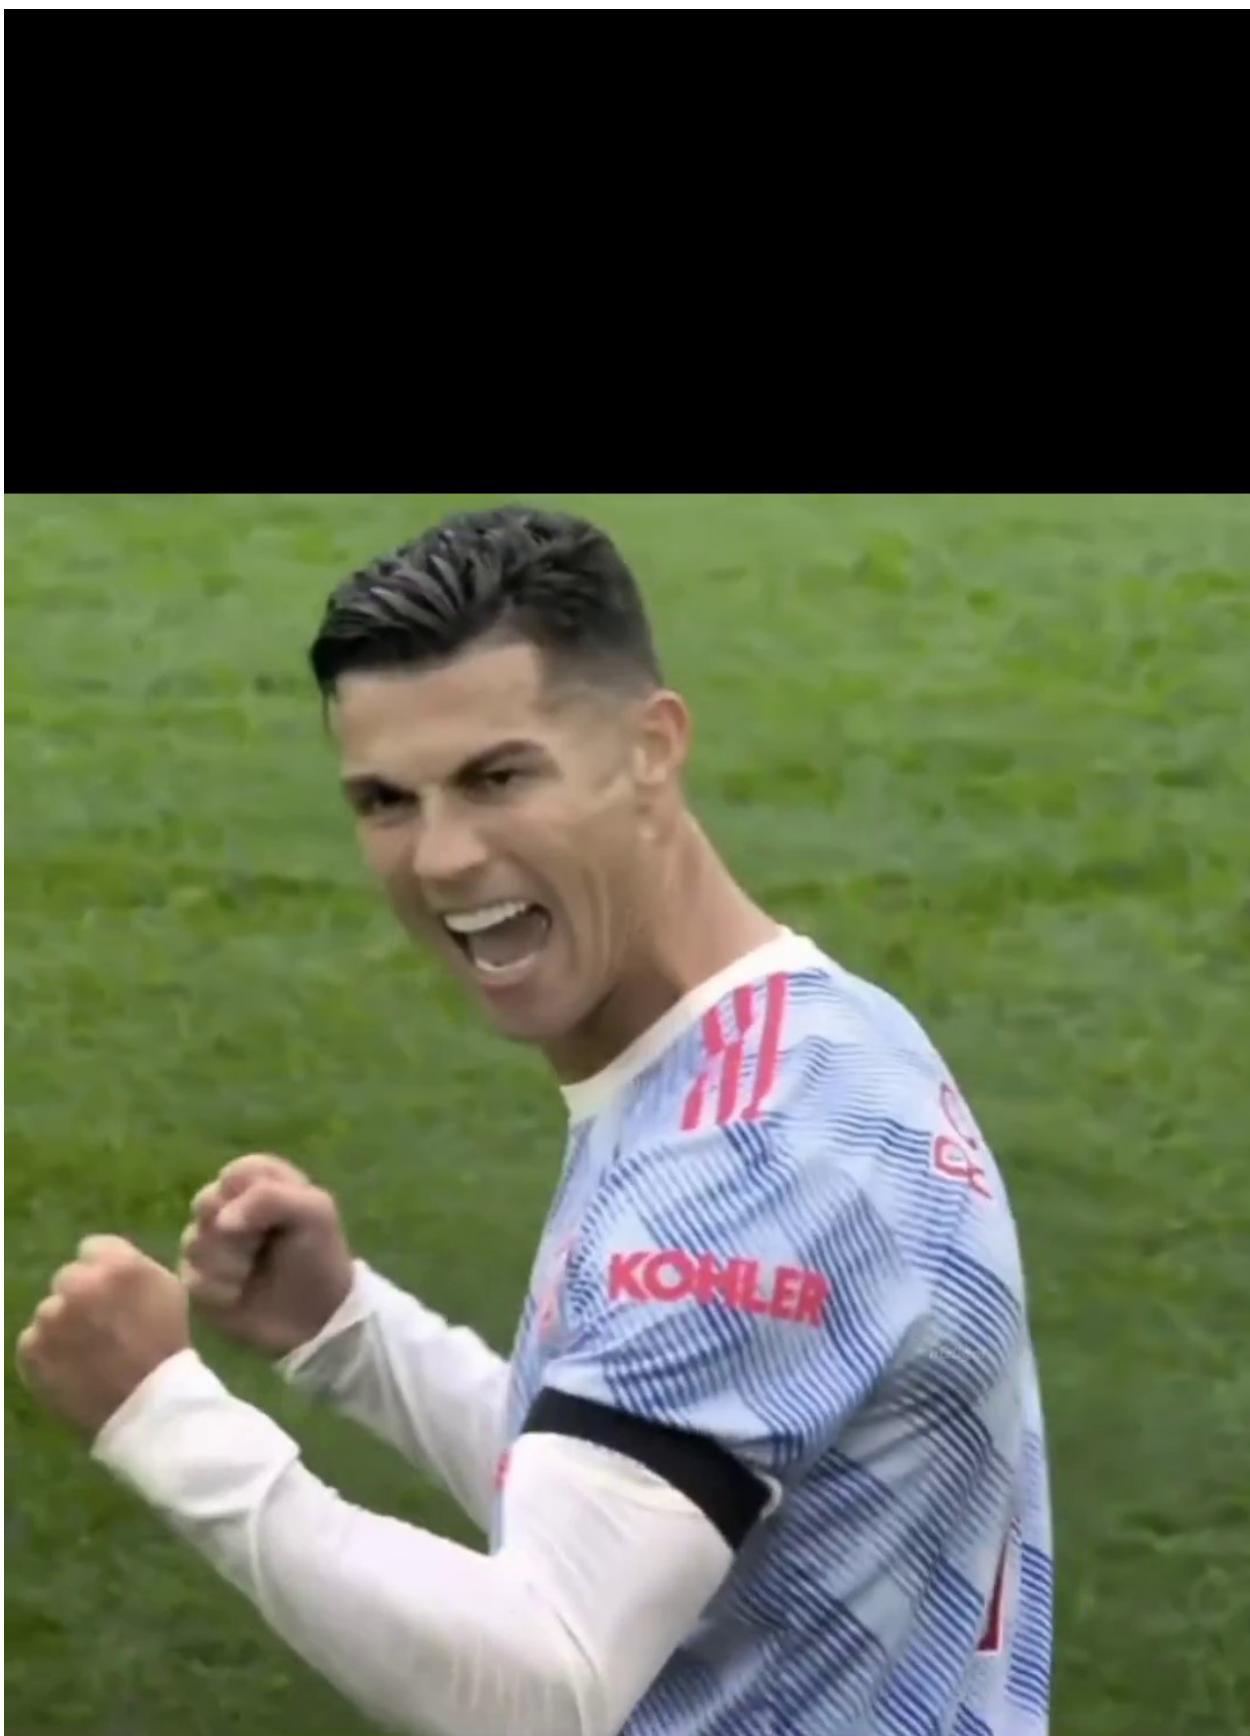
0.15366512345678984

Change detected



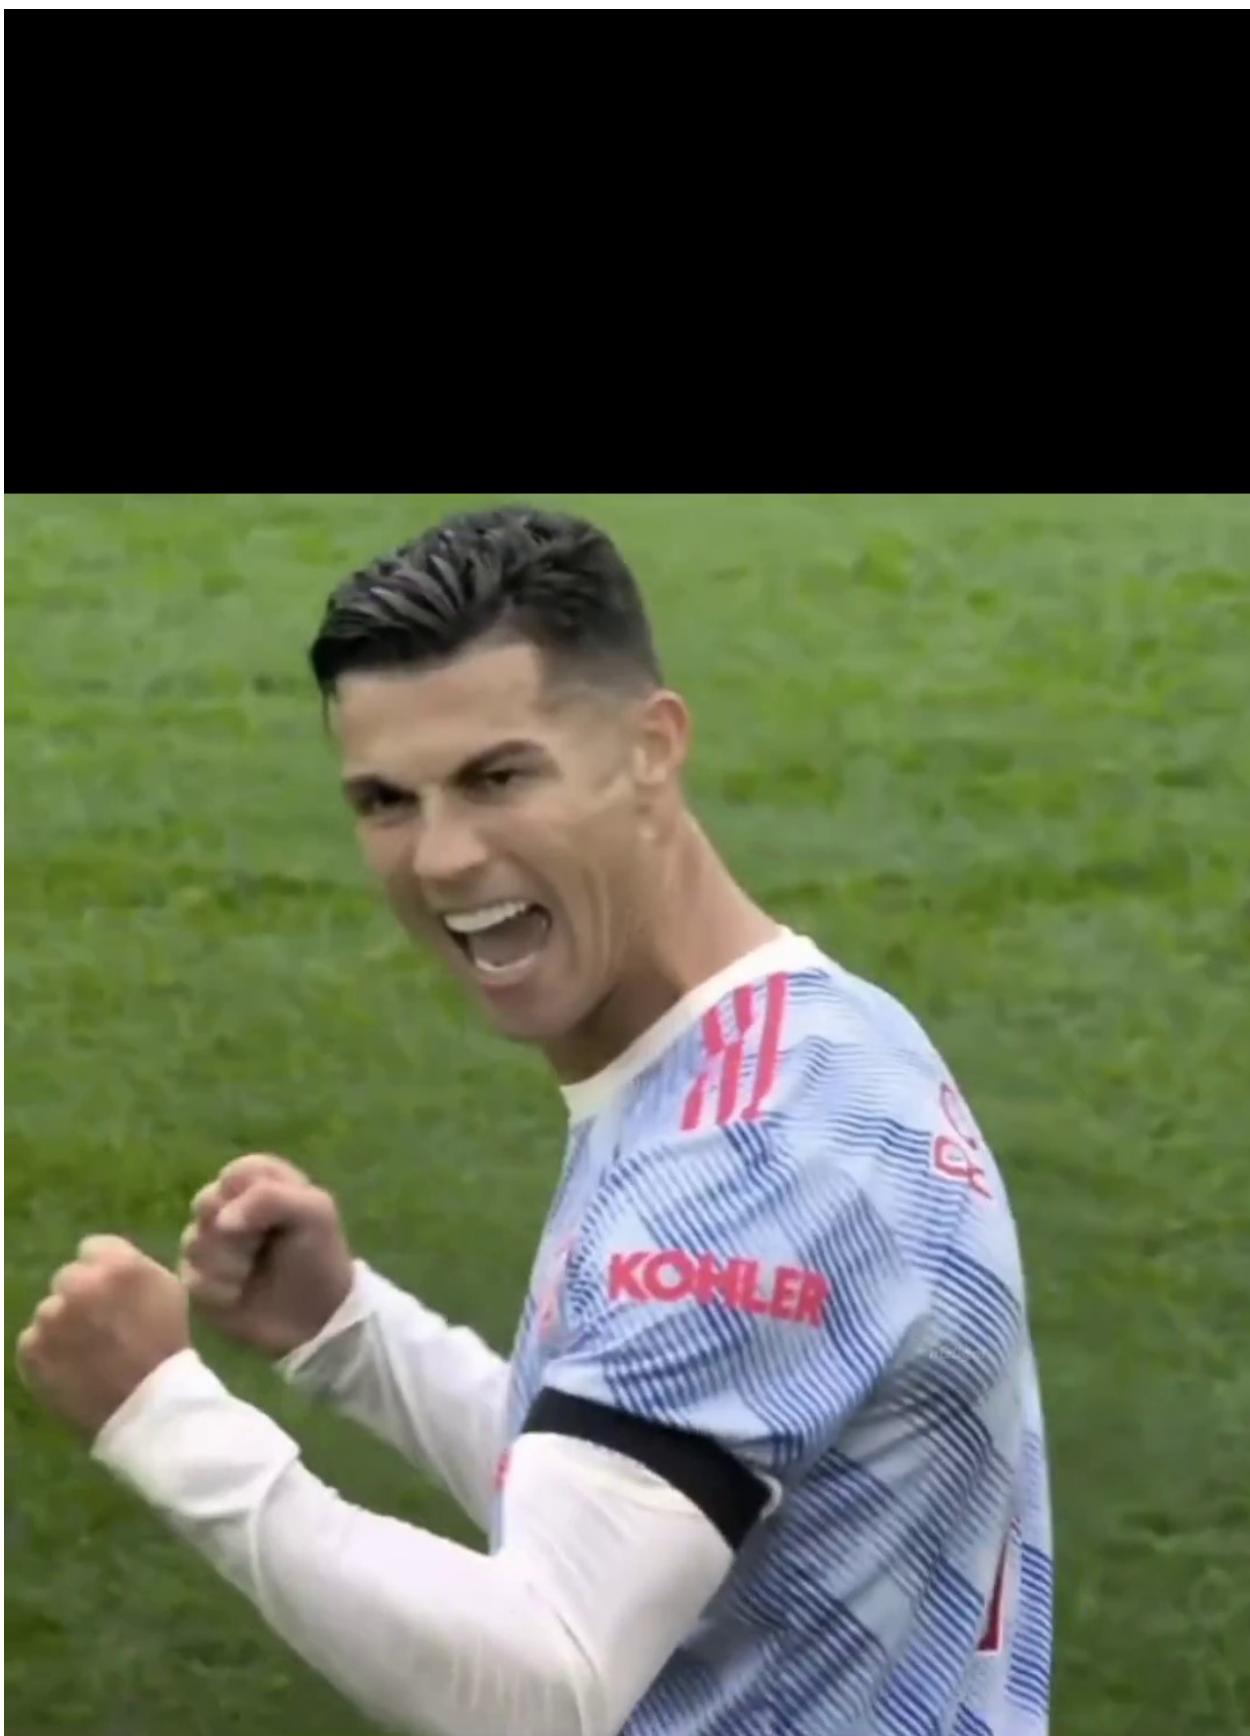
0.08837191358024654

Change detected



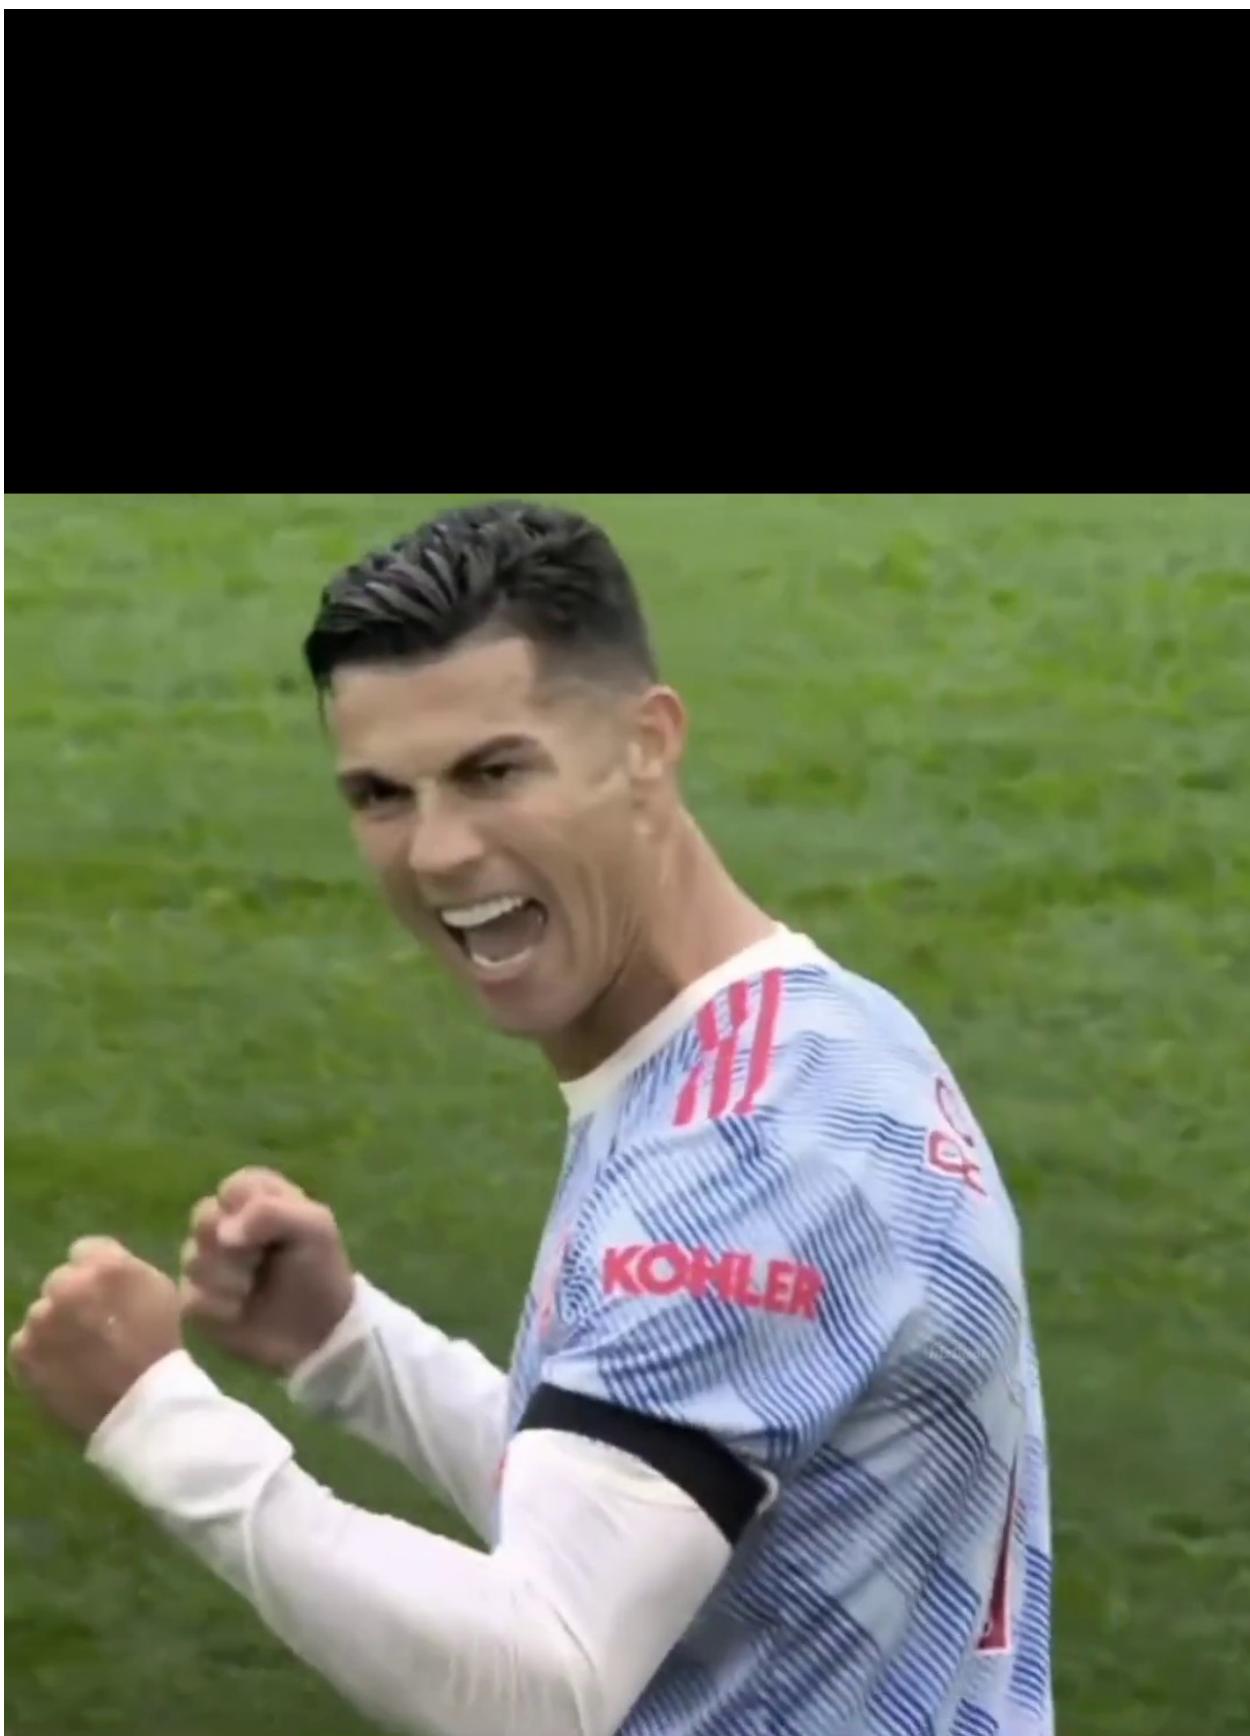
0.006894290123454994

Change detected



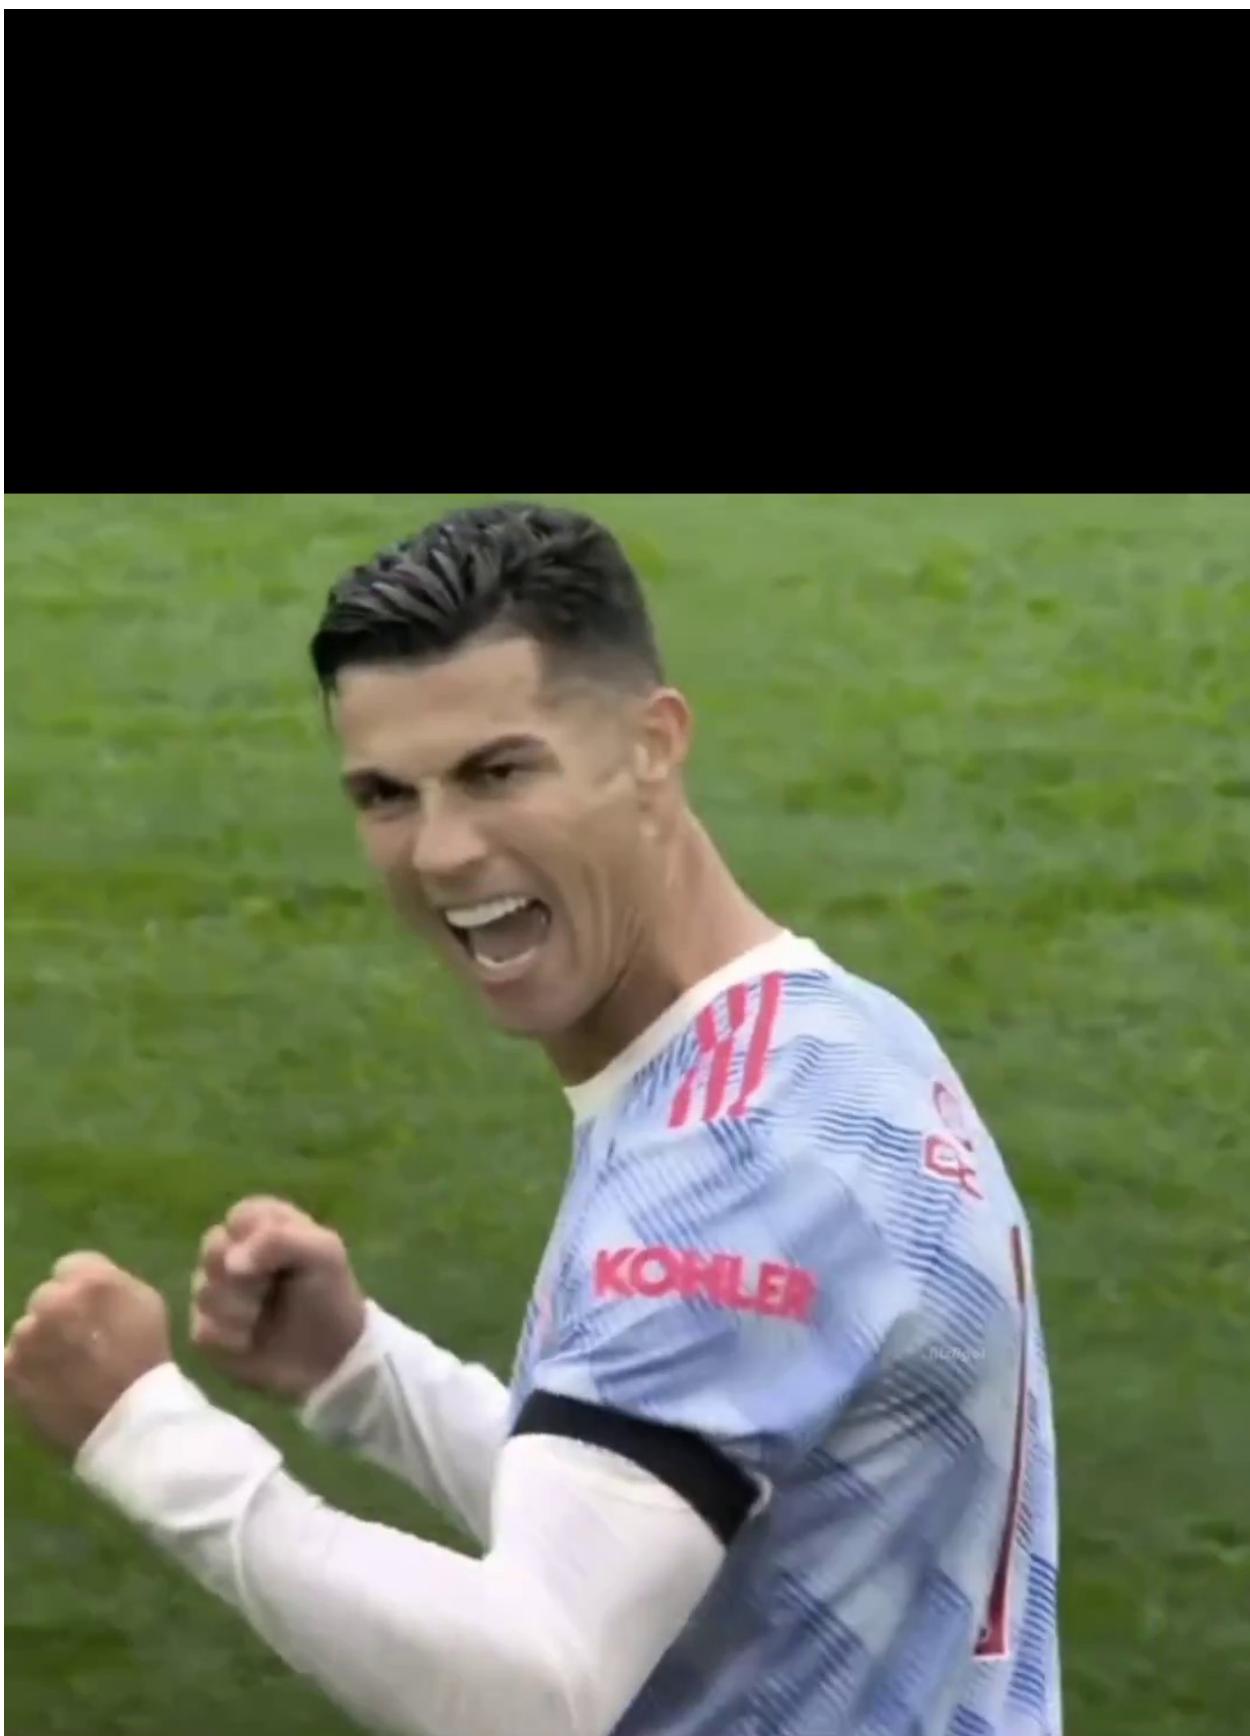
0.03742717978394694

Change detected



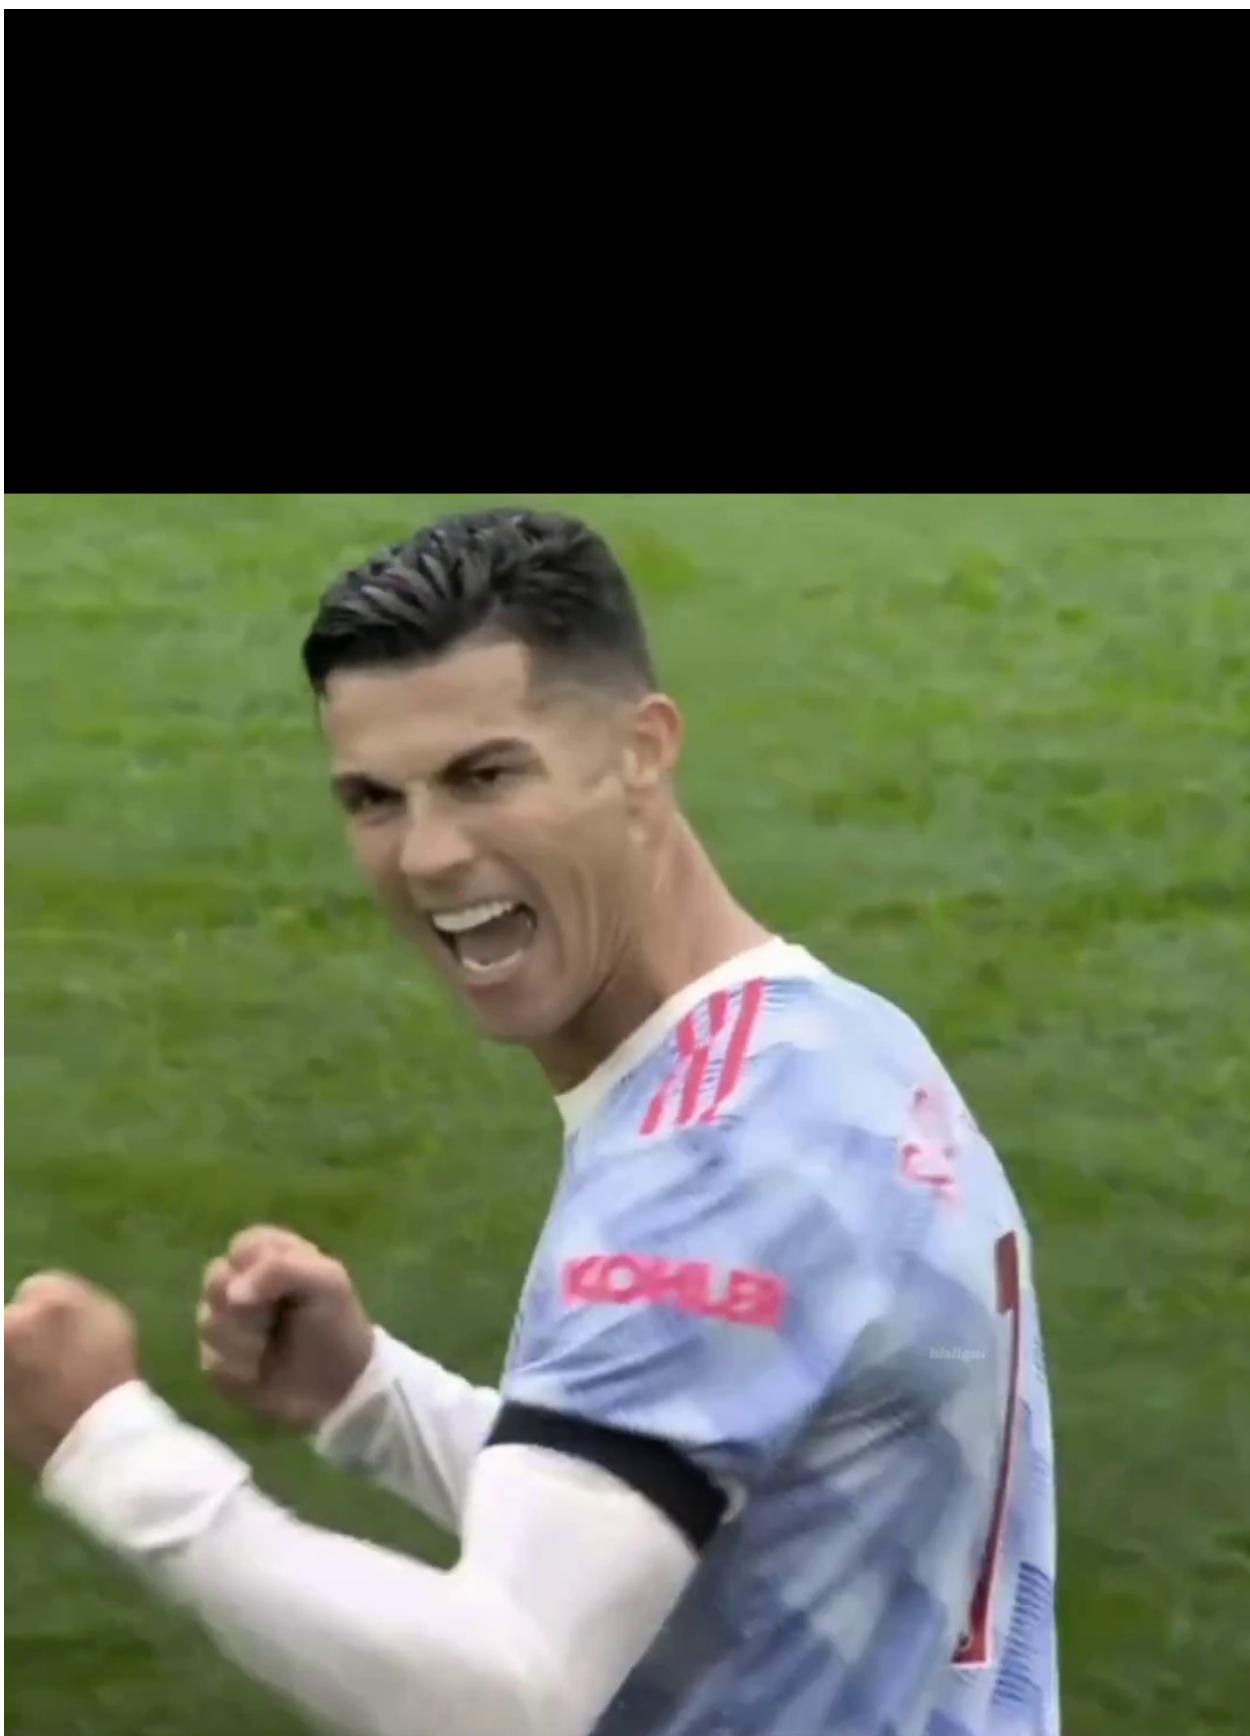
0.10041087962963502

Change detected



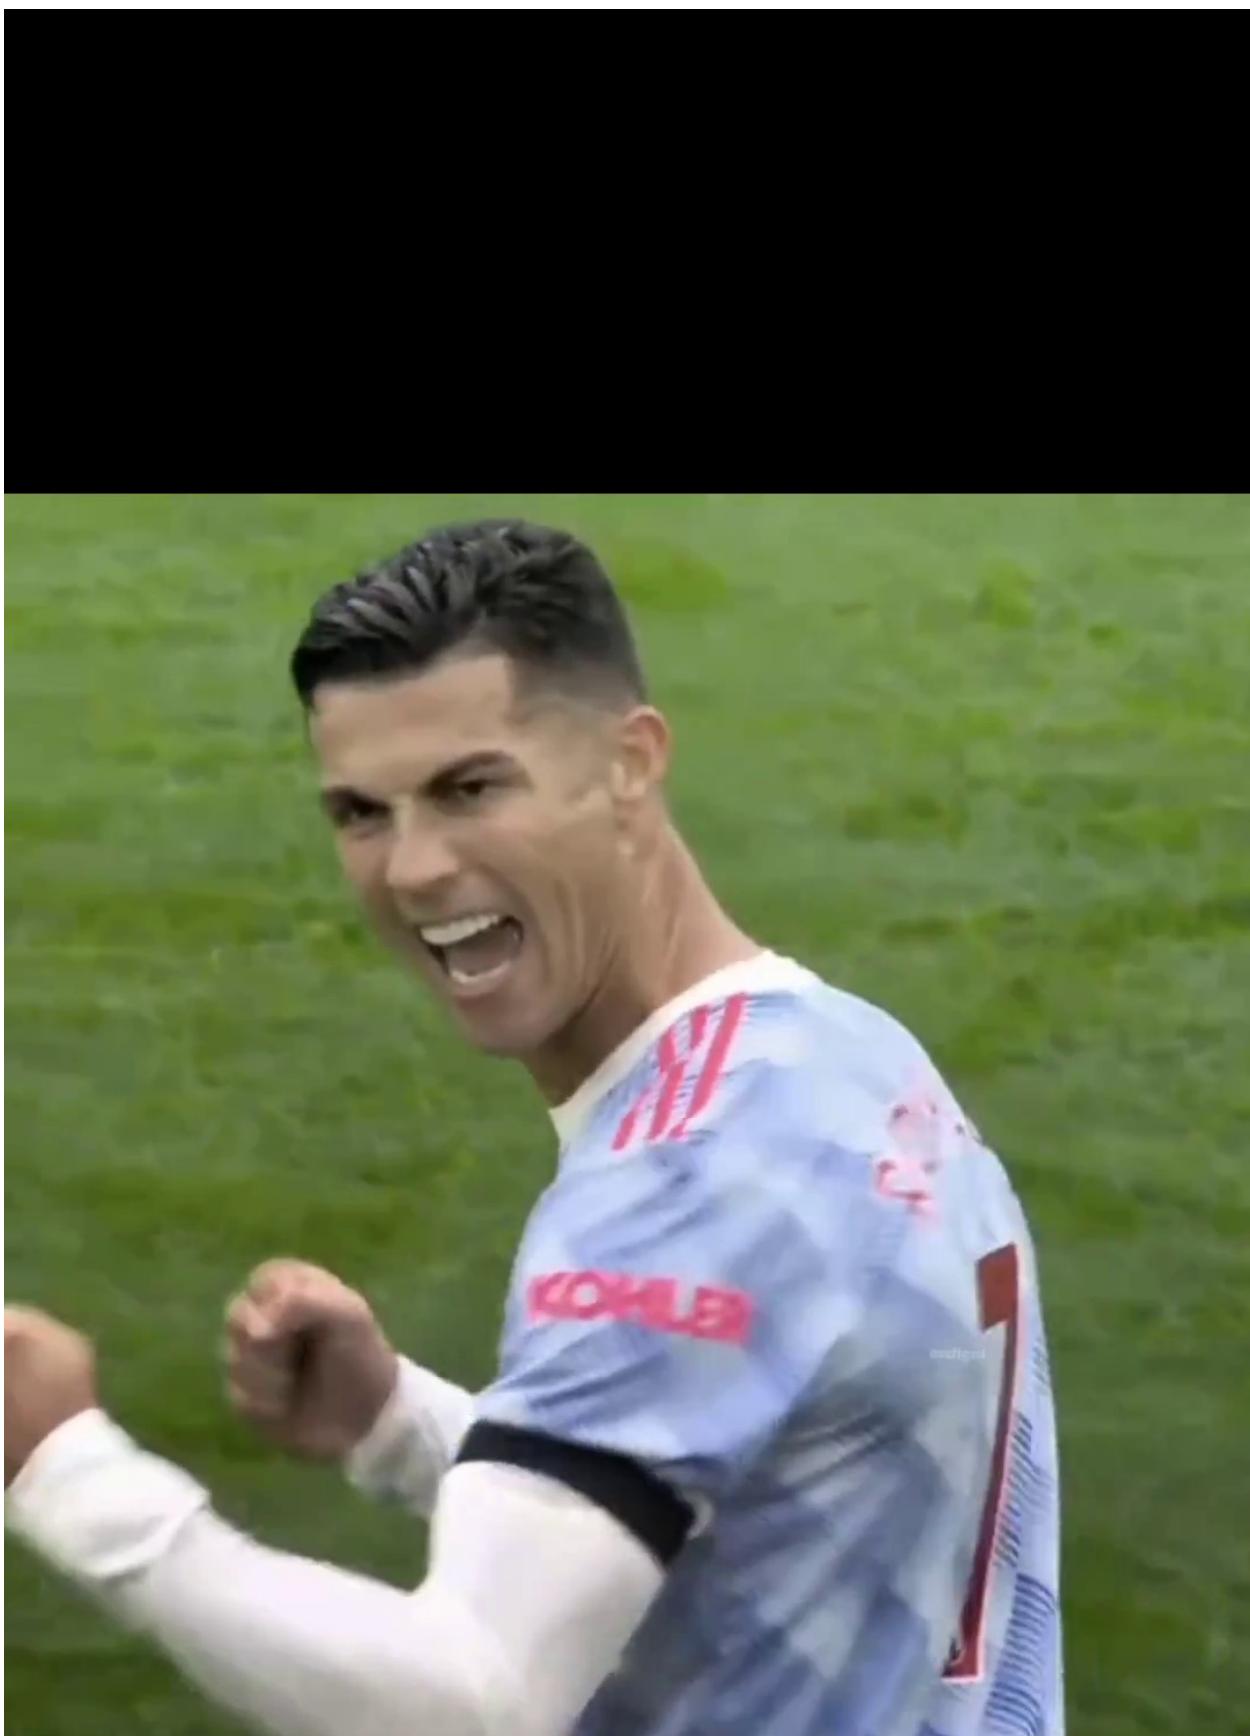
0.22392216435184764

Change detected



0.37755642361112507

Change detected



0.10220920138890222

Change detected



19.574018614969127

Change detected



hizligol

0.24344232253086773

Change detected



hizligol

0.022759934413578264

Change detected



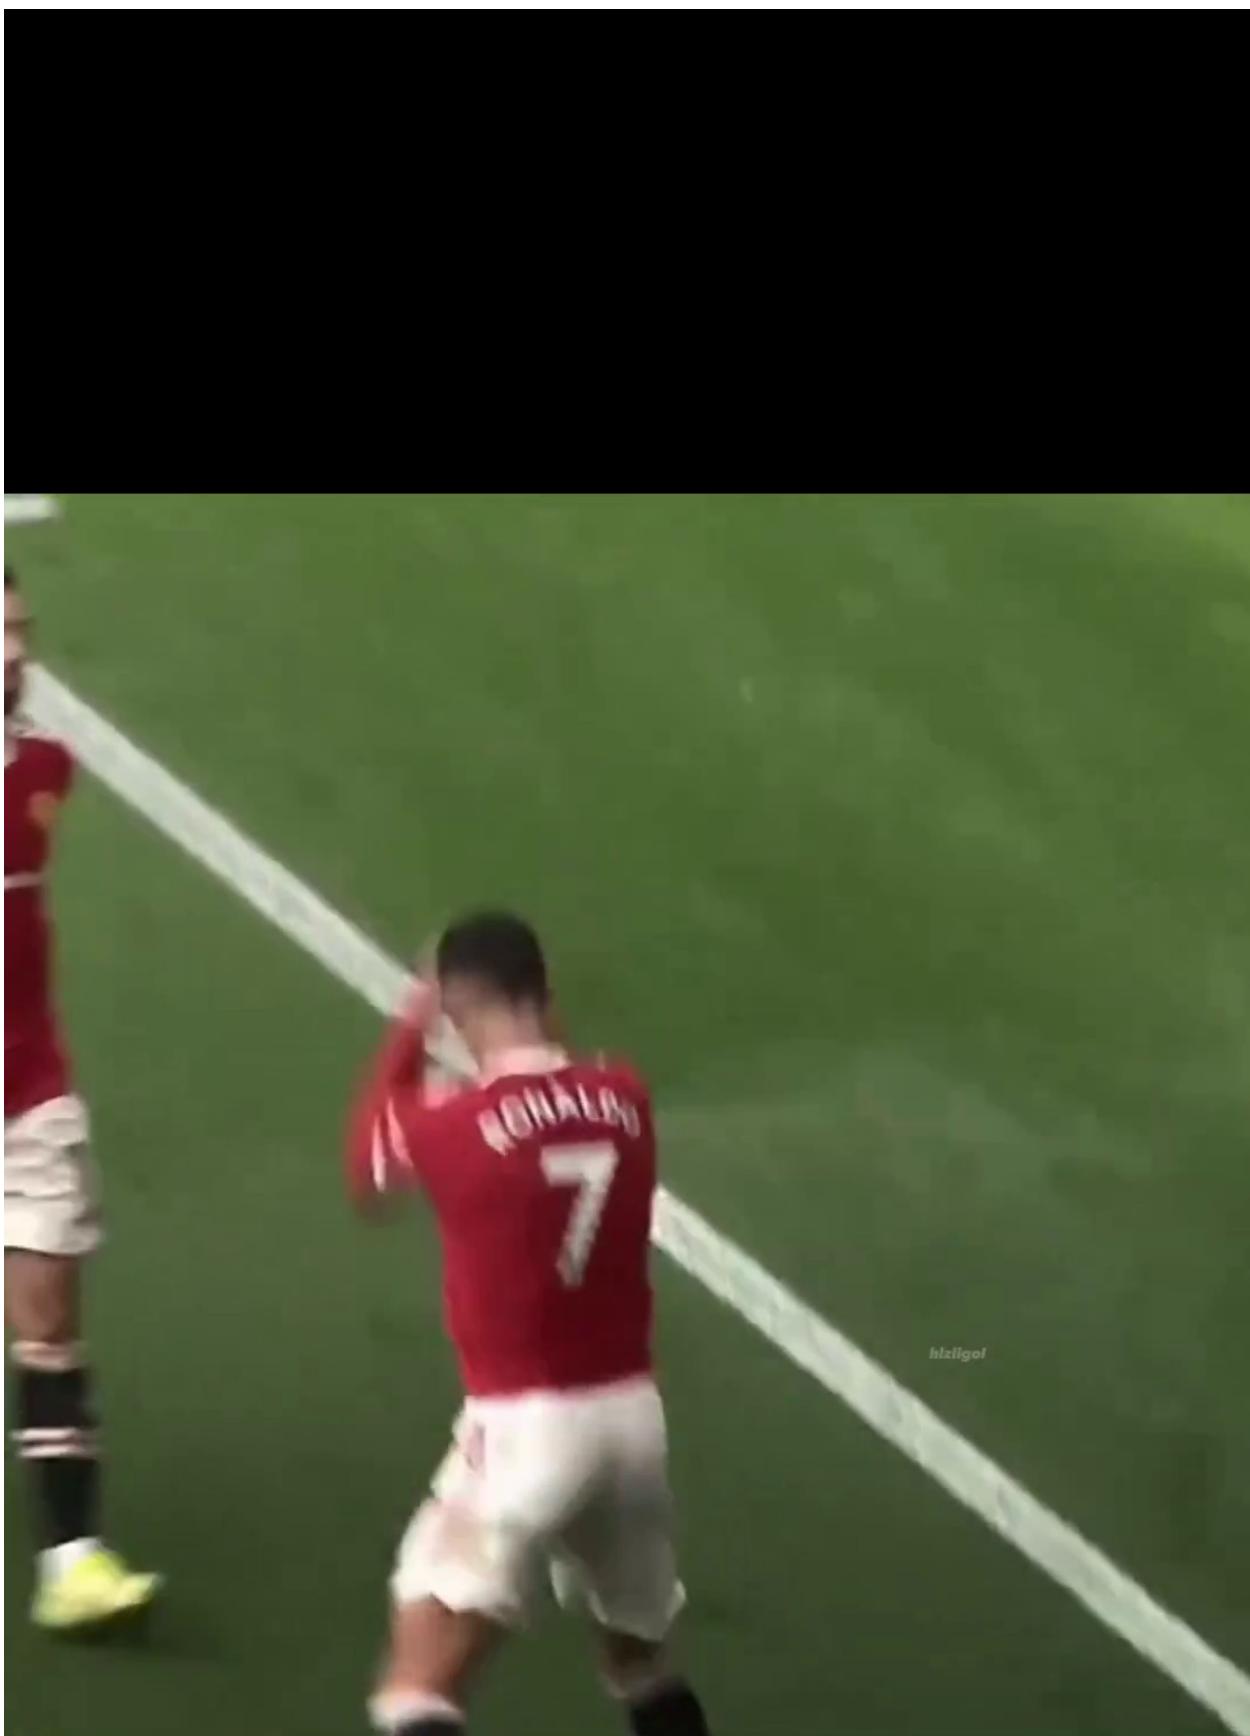
0.013610146604939644

Change detected



0.5110966435185134

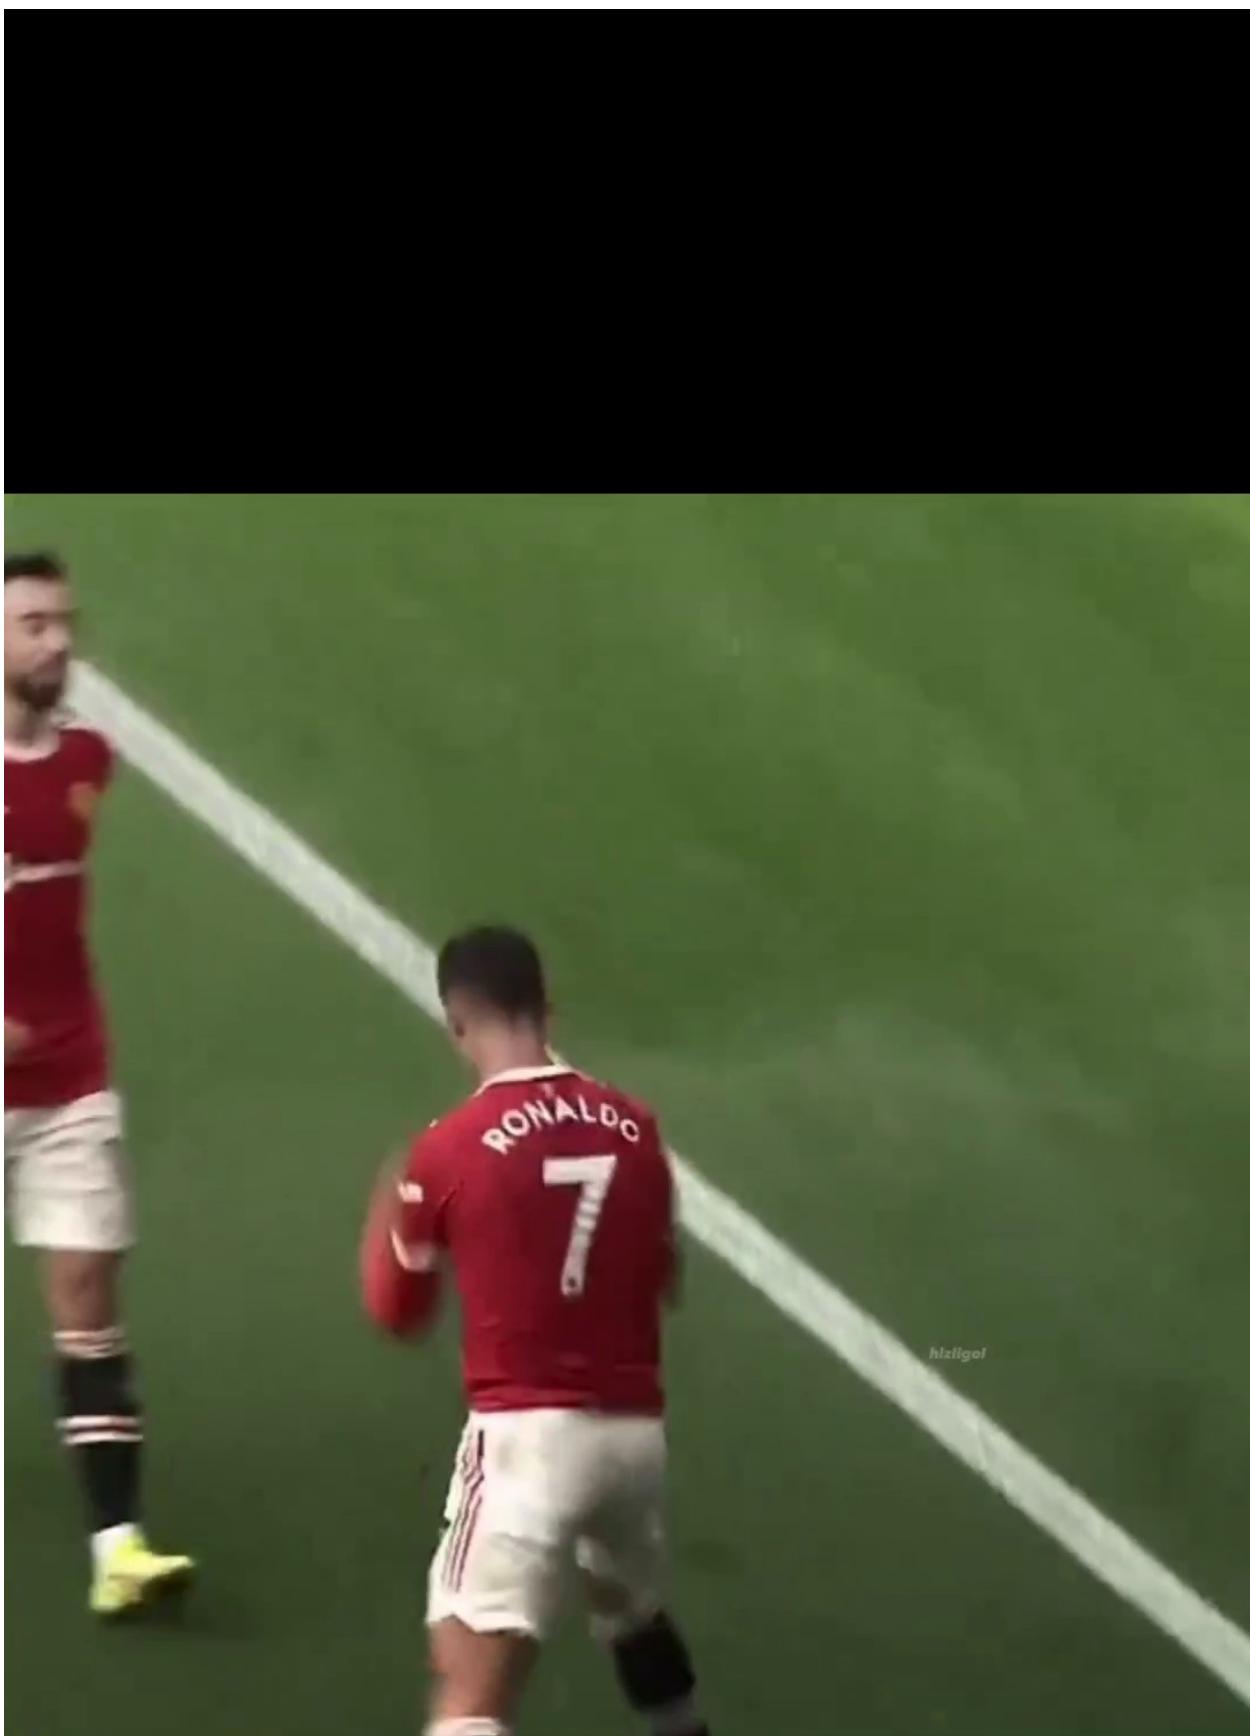
Change detected



hizligol

0.15378327546296333

Change detected



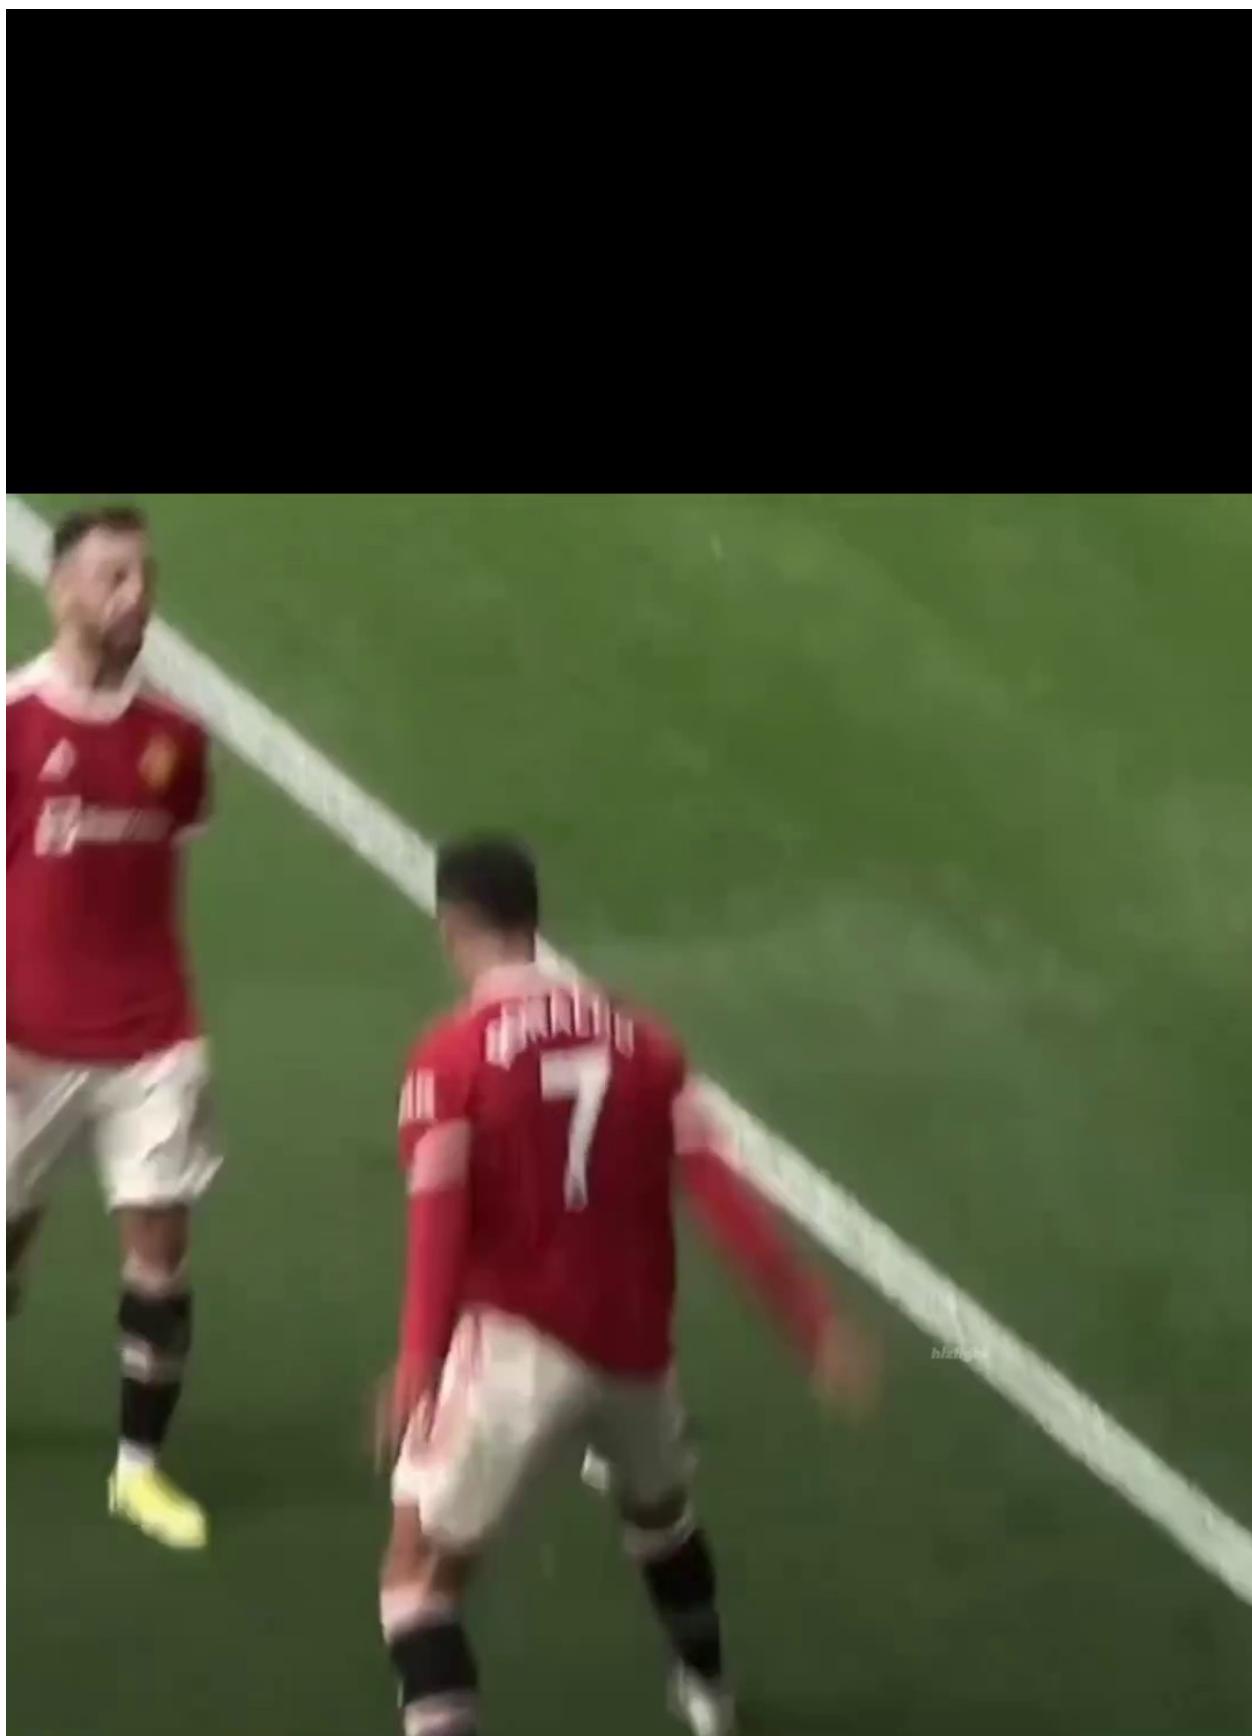
0.8563551311728332

Change detected



0.14303915895062147

Change detected



0.288315972222217

Change detected



0.02205825617284063

Change detected



0.12366560570987417

Change detected



0.16315923996913995

Change detected



hizligol

0.2567496141975312

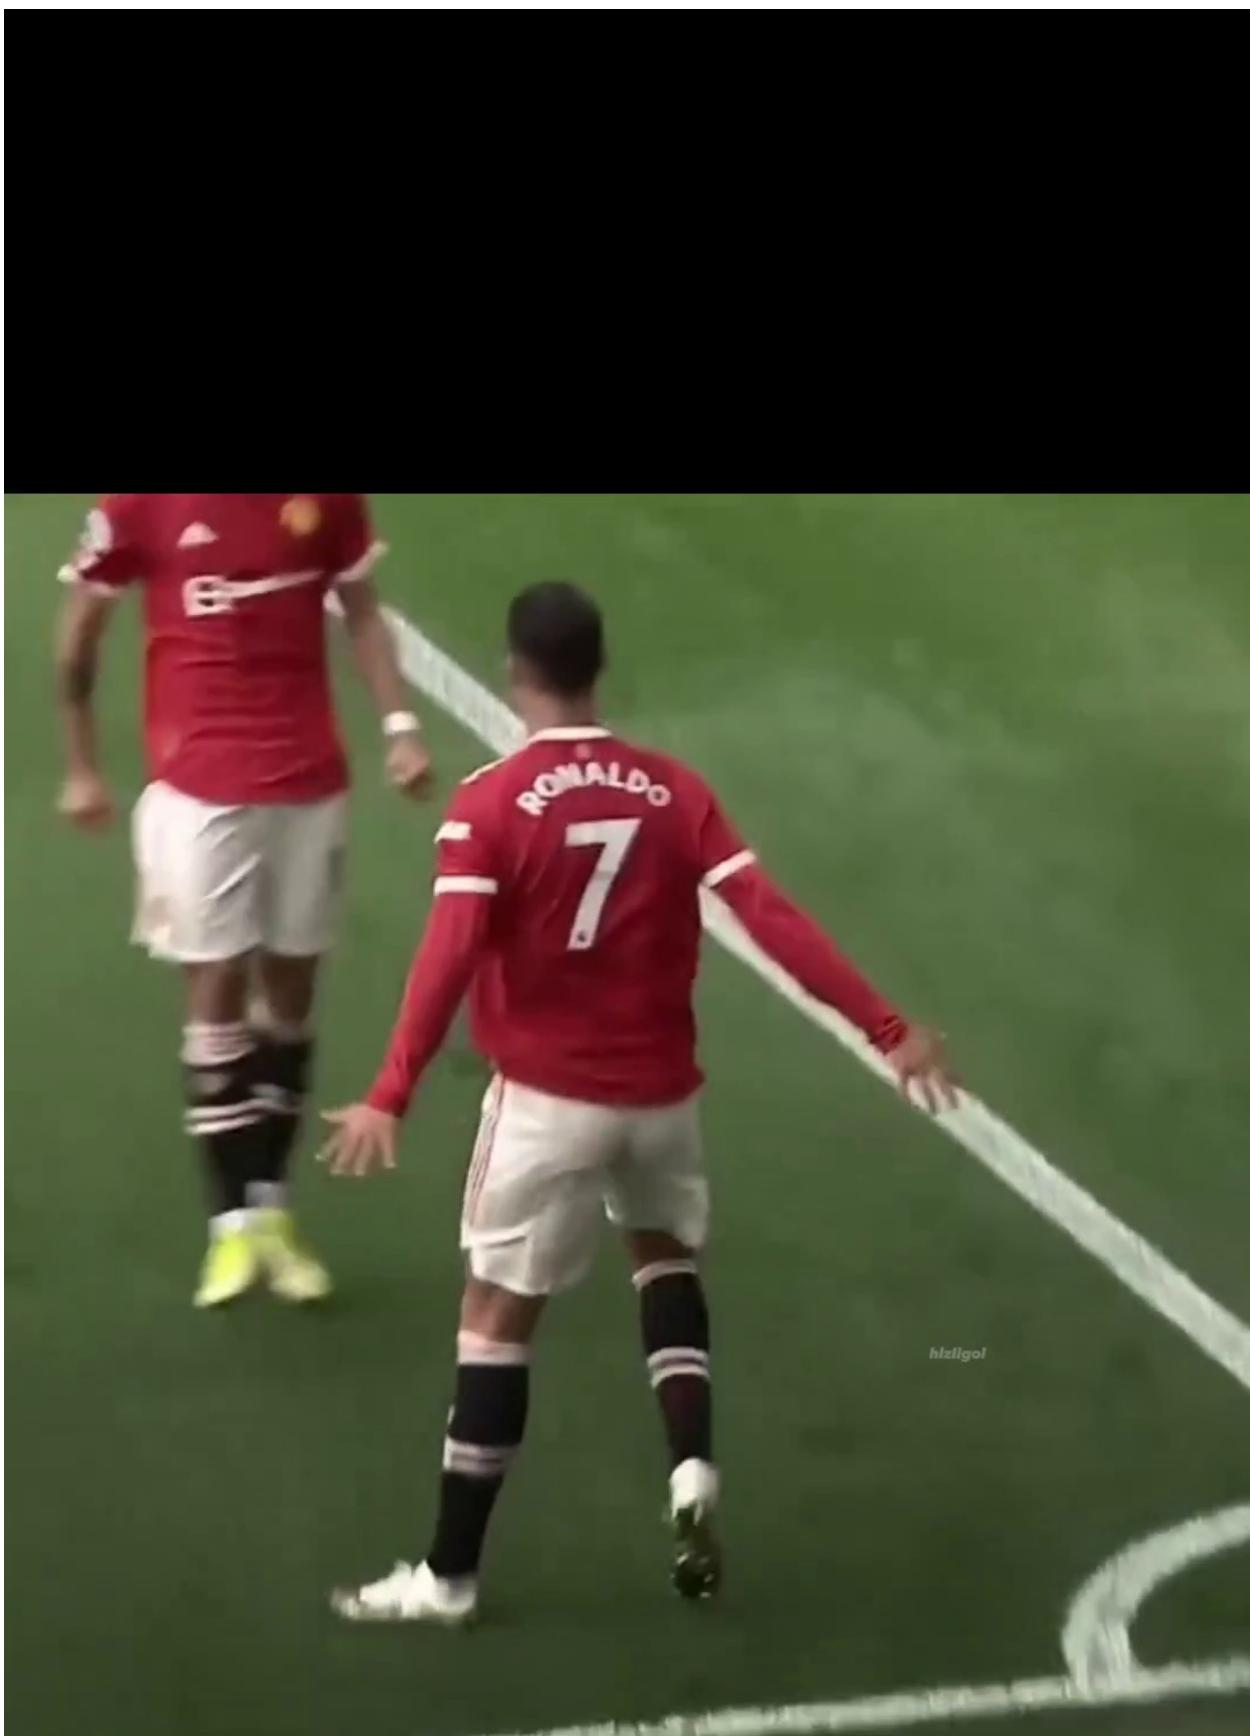
Change detected



hizligol

0.21685619212962592

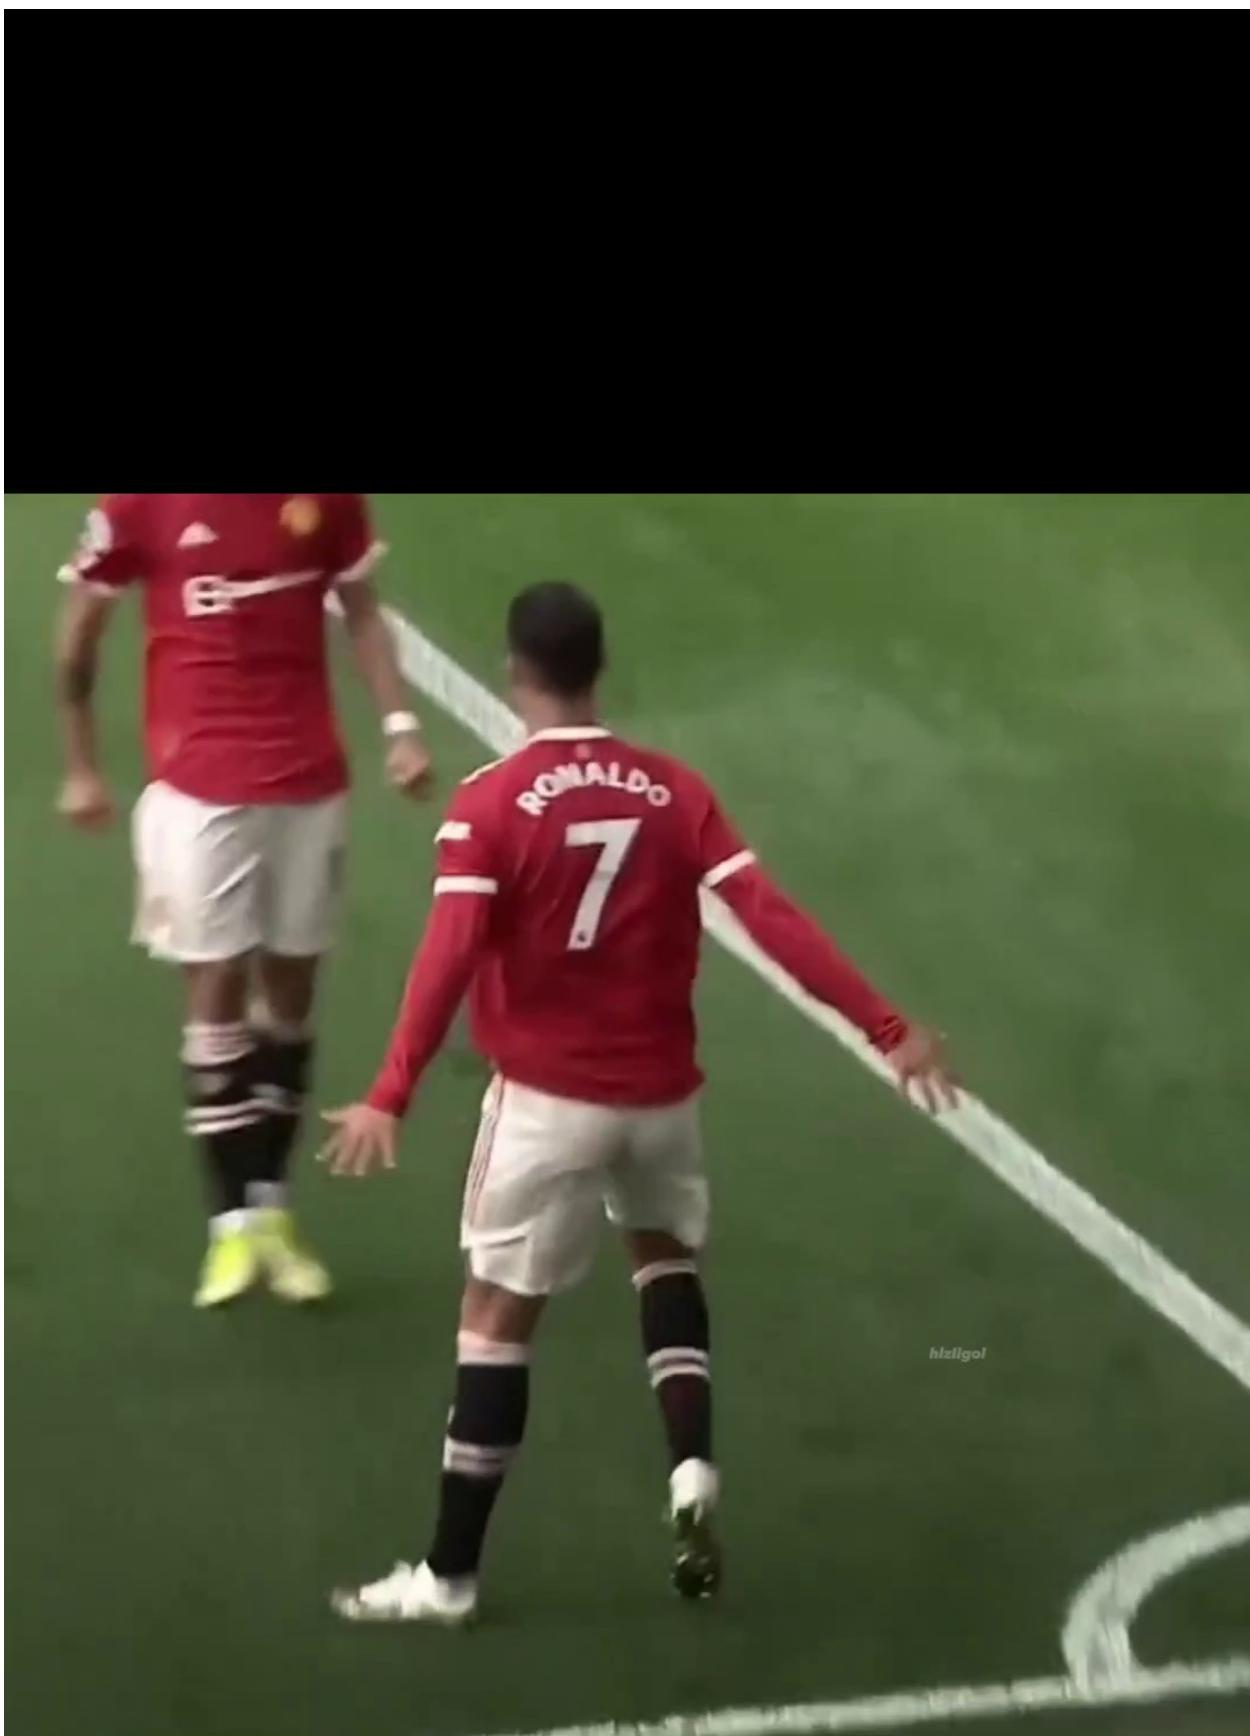
Change detected



hizligol

0.10874710648148067

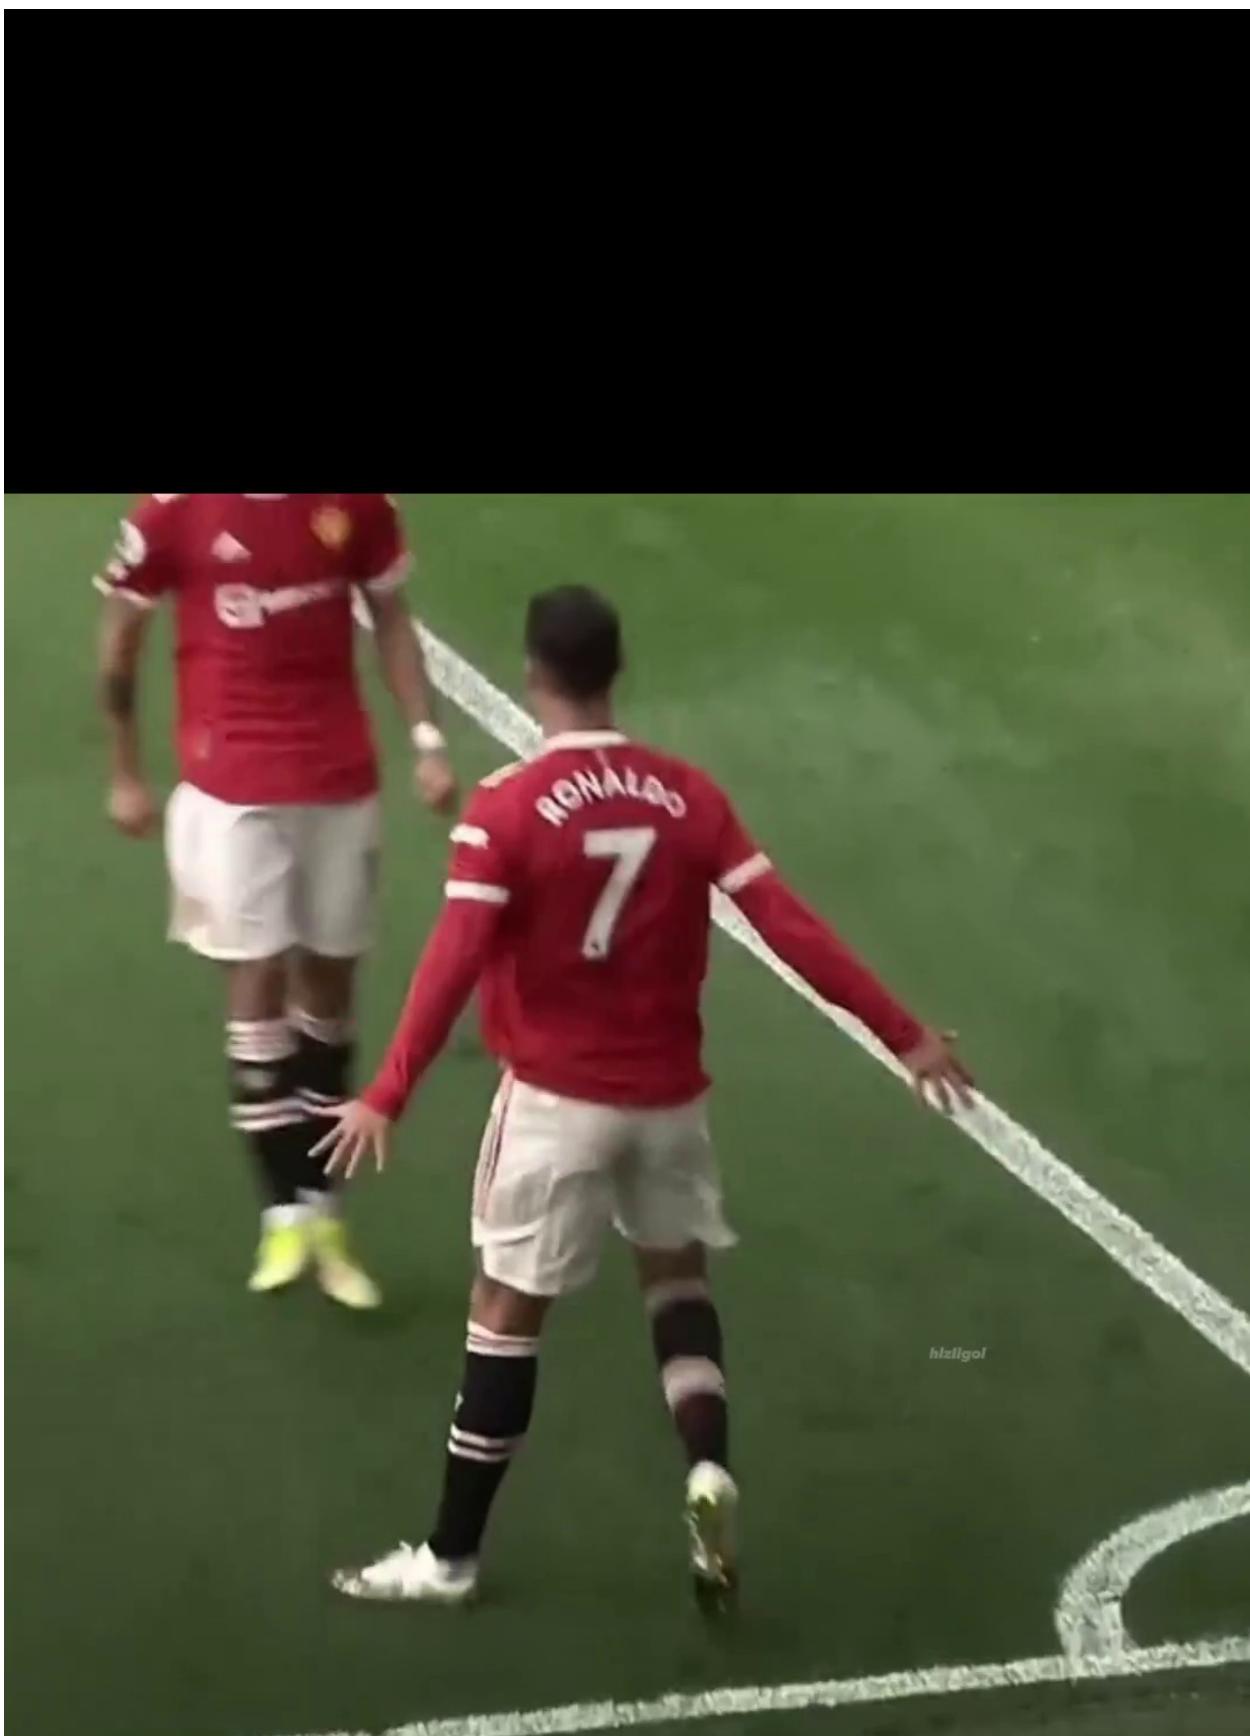
Change detected



hizligol

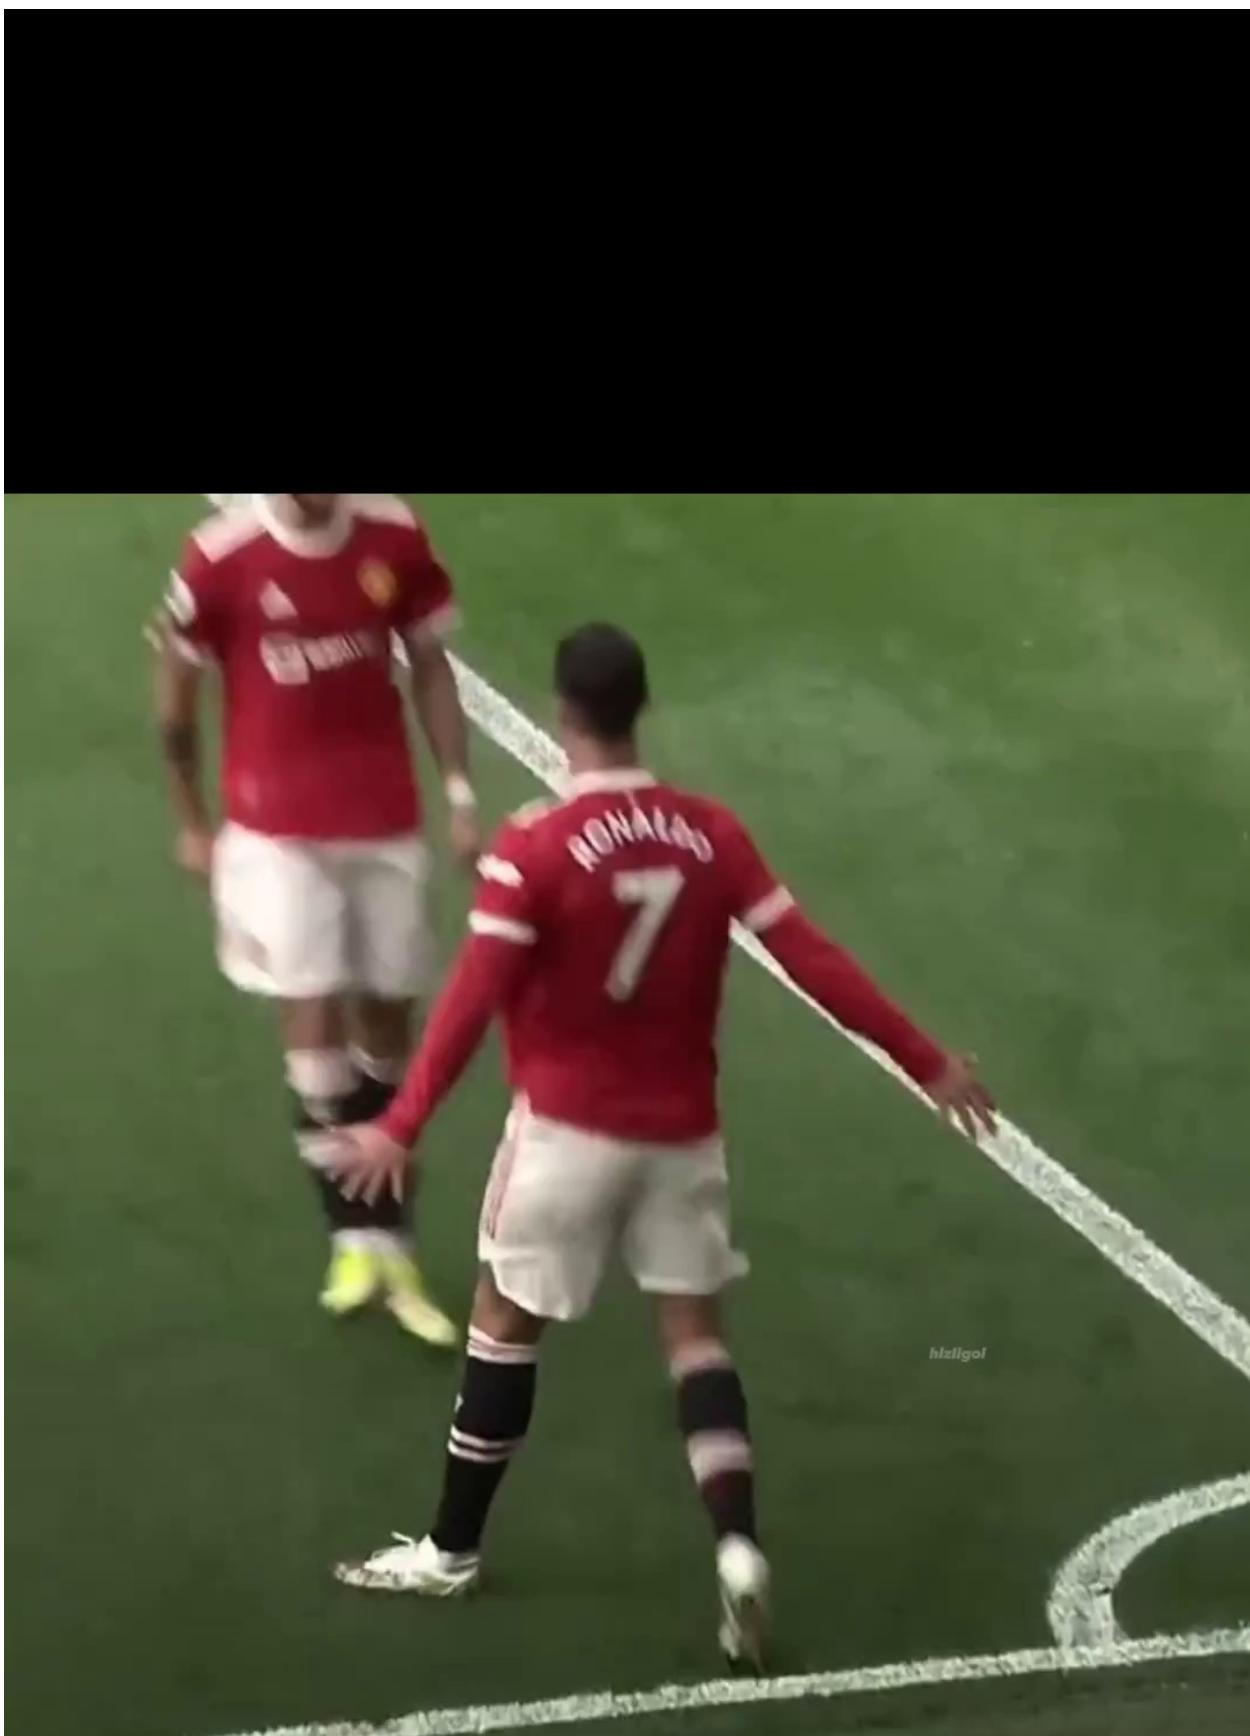
0.0020351080246925335

Change detected



0.07265383873456699

Change detected



hizligol

0.5752131558641977

Change detected



0.08875385802468827

Change detected



hizligol

```
#Practical-10-Programs for Face Detection
```

```
pip install opencv-python-headless matplotlib

Requirement already satisfied: opencv-python-headless in
/usr/local/lib/python3.10/dist-packages (4.10.0.82)
Requirement already satisfied: matplotlib in
/usr/local/lib/python3.10/dist-packages (3.7.1)
Requirement already satisfied: numpy>=1.21.2 in
/usr/local/lib/python3.10/dist-packages (from opencv-python-headless)
(1.25.2)
Requirement already satisfied: contourpy>=1.0.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (1.2.1)
Requirement already satisfied: cycler>=0.10 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (4.53.0)
Requirement already satisfied: kiwisolver>=1.0.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (1.4.5)
Requirement already satisfied: packaging>=20.0 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (24.1)
Requirement already satisfied: pillow>=6.2.0 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (9.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (3.1.2)
Requirement already satisfied: python-dateutil>=2.7 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (2.8.2)
Requirement already satisfied: six>=1.5 in
/usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7-
>matplotlib) (1.16.0)

import urllib.request

# URL of the Haar Cascade file
url =
'https://github.com/opencv/opencv/raw/master/data/haarcascades/haarcascade_frontalface_default.xml'

# Path where the file will be saved
haar_cascade_path = 'haarcascade_frontalface_default.xml'

# Download the file
urllib.request.urlretrieve(url, haar_cascade_path)

('haarcascade_frontalface_default.xml',
 <http.client.HTTPMessage at 0x7e1dbd13a5c0>)

import os
import numpy as np
import cv2
import matplotlib.pyplot as plt
```

```

import urllib.request

# Download Haar Cascade if not already present
haar_cascade_path = 'haarcascade_frontalface_default.xml'
if not os.path.exists(haar_cascade_path):
    url =
'https://github.com/opencv/opencv/raw/master/data/haarcascades/haarcascade_frontalface_default.xml'
    urllib.request.urlretrieve(url, haar_cascade_path)

# Define the image path
img_path = 'Galaxy11.jpg'

# Load the image
img = cv2.imread(img_path)

# Check if the image was loaded properly
if img is None:
    print(f"Error: Could not load image from {img_path}")
else:
    # Convert the image from BGR to RGB
    img_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

    # Display the original image
    plt.figure(figsize=(8, 6))
    plt.imshow(img_rgb)
    plt.axis('off')
    plt.title('Original Image')
    plt.show()

    # Convert the image to grayscale
    gray_img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

# Load the Haar Cascade
haar_cascade = cv2.CascadeClassifier(haar_cascade_path)

# Check if the Haar Cascade was loaded properly
if haar_cascade.empty():
    print(f"Error: Could not load Haar Cascade from {haar_cascade_path}")
else:
    # Detect faces
    faces_rect = haar_cascade.detectMultiScale(gray_img,
scaleFactor=1.1, minNeighbors=9)

    # Draw rectangles around detected faces
    for (x, y, w, h) in faces_rect:
        cv2.rectangle(img_rgb, (x, y), (x + w, y + h), (0, 255, 0), 2)

```

```
# Display the image with detected faces
plt.figure(figsize=(8, 6))
plt.imshow(img_rgb)
plt.axis('off')
plt.title('Image with Detected Faces')
plt.show()
```

Original Image



Image with Detected Faces



```

#Practical-11_Object_Detection
#Note-Perform it in Google Collab because the belloe libraries has no
support in Jupyter Notebook

# Step 1: Install Required Libraries
!pip install ultralytics
!pip install matplotlib
!pip install opencv-python-headless

# Step 2: Import Libraries and Load the Model
import matplotlib.pyplot as plt
from ultralytics import YOLO
import cv2

# Load the YOLOv8 model
model = YOLO('yolov8n.pt')

# Step 3: Perform Object Detection and Save Results
# Perform object detection on the input image
image_path = '/content/Emue.jpg'
results = model(image_path)

# Iterate over the results and extract information
for i, result in enumerate(results):
    # Display and save each result
    result.show()
    result_image_path = f'result_{i}.jpg'
    result.save(result_image_path)

# Step 4: Display Original and Result Image Side by Side
# Load the original image
original_image = cv2.imread(image_path)
original_image_rgb = cv2.cvtColor(original_image, cv2.COLOR_BGR2RGB)

# Load and display all result images
plt.figure(figsize=(20, 10))

plt.subplot(1, 2, 1)
plt.imshow(original_image_rgb)
plt.axis('on')
plt.title('Original Image')

for i in range(len(results)):
    result_image = cv2.imread(f'result_{i}.jpg')
    result_image_rgb = cv2.cvtColor(result_image, cv2.COLOR_BGR2RGB)
    plt.subplot(1, 2, 2)
    plt.imshow(result_image_rgb)
    plt.axis('on')
    plt.title('YOLOv8 Object Detection Result')

```

```
plt.show()

Requirement already satisfied: ultralytics in
/usr/local/lib/python3.10/dist-packages (8.2.35)
Requirement already satisfied: numpy<2.0.0 in
/usr/local/lib/python3.10/dist-packages (from ultralytics) (1.25.2)
Requirement already satisfied: matplotlib>=3.3.0 in
/usr/local/lib/python3.10/dist-packages (from ultralytics) (3.7.1)
Requirement already satisfied: opencv-python>=4.6.0 in
/usr/local/lib/python3.10/dist-packages (from ultralytics) (4.8.0.76)
Requirement already satisfied: pillow>=7.1.2 in
/usr/local/lib/python3.10/dist-packages (from ultralytics) (9.4.0)
Requirement already satisfied: pyyaml>=5.3.1 in
/usr/local/lib/python3.10/dist-packages (from ultralytics) (6.0.1)
Requirement already satisfied: requests>=2.23.0 in
/usr/local/lib/python3.10/dist-packages (from ultralytics) (2.31.0)
Requirement already satisfied: scipy>=1.4.1 in
/usr/local/lib/python3.10/dist-packages (from ultralytics) (1.11.4)
Requirement already satisfied: torch>=1.8.0 in
/usr/local/lib/python3.10/dist-packages (from ultralytics)
(2.3.0+cu121)
Requirement already satisfied: torchvision>=0.9.0 in
/usr/local/lib/python3.10/dist-packages (from ultralytics)
(0.18.0+cu121)
Requirement already satisfied: tqdm>=4.64.0 in
/usr/local/lib/python3.10/dist-packages (from ultralytics) (4.66.4)
Requirement already satisfied: psutil in
/usr/local/lib/python3.10/dist-packages (from ultralytics) (5.9.5)
Requirement already satisfied: py-cpuinfo in
/usr/local/lib/python3.10/dist-packages (from ultralytics) (9.0.0)
Requirement already satisfied: pandas>=1.1.4 in
/usr/local/lib/python3.10/dist-packages (from ultralytics) (2.0.3)
Requirement already satisfied: seaborn>=0.11.0 in
/usr/local/lib/python3.10/dist-packages (from ultralytics) (0.13.1)
Requirement already satisfied: ultralytics-thop>=2.0.0 in
/usr/local/lib/python3.10/dist-packages (from ultralytics) (2.0.0)
Requirement already satisfied: contourpy>=1.0.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib>=3.3.0-
>ultralytics) (1.2.1)
Requirement already satisfied: cycler>=0.10 in
/usr/local/lib/python3.10/dist-packages (from matplotlib>=3.3.0-
>ultralytics) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in
/usr/local/lib/python3.10/dist-packages (from matplotlib>=3.3.0-
>ultralytics) (4.53.0)
Requirement already satisfied: kiwisolver>=1.0.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib>=3.3.0-
>ultralytics) (1.4.5)
Requirement already satisfied: packaging>=20.0 in
```

```
/usr/local/lib/python3.10/dist-packages (from matplotlib>=3.3.0->ultralytics) (24.1)
Requirement already satisfied: pyparsing>=2.3.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib>=3.3.0->ultralytics) (3.1.2)
Requirement already satisfied: python-dateutil>=2.7 in
/usr/local/lib/python3.10/dist-packages (from matplotlib>=3.3.0->ultralytics) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in
/usr/local/lib/python3.10/dist-packages (from pandas>=1.1.4->ultralytics) (2023.4)
Requirement already satisfied: tzdata>=2022.1 in
/usr/local/lib/python3.10/dist-packages (from pandas>=1.1.4->ultralytics) (2024.1)
Requirement already satisfied: charset-normalizer<4,>=2 in
/usr/local/lib/python3.10/dist-packages (from requests>=2.23.0->ultralytics) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in
/usr/local/lib/python3.10/dist-packages (from requests>=2.23.0->ultralytics) (3.7)
Requirement already satisfied: urllib3<3,>=1.21.1 in
/usr/local/lib/python3.10/dist-packages (from requests>=2.23.0->ultralytics) (2.0.7)
Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.10/dist-packages (from requests>=2.23.0->ultralytics) (2024.6.2)
Requirement already satisfied: filelock in
/usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (3.14.0)
Requirement already satisfied: typing-extensions>=4.8.0 in
/usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (4.12.2)
Requirement already satisfied: sympy in
/usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (1.12.1)
Requirement already satisfied: networkx in
/usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (3.3)
Requirement already satisfied: jinja2 in
/usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (3.1.4)
Requirement already satisfied: fsspec in
/usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (2023.6.0)
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.1.105 in
/usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (12.1.105)
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.1.105
in /usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (12.1.105)
```

```
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.1.105 in
/usr/local/lib/python3.10/dist-packages (from torch>=1.8.0-
>ultralytics) (12.1.105)
Requirement already satisfied: nvidia-cudnn-cu12==8.9.2.26 in
/usr/local/lib/python3.10/dist-packages (from torch>=1.8.0-
>ultralytics) (8.9.2.26)
Requirement already satisfied: nvidia-cublas-cu12==12.1.3.1 in
/usr/local/lib/python3.10/dist-packages (from torch>=1.8.0-
>ultralytics) (12.1.3.1)
Requirement already satisfied: nvidia-cufft-cu12==11.0.2.54 in
/usr/local/lib/python3.10/dist-packages (from torch>=1.8.0-
>ultralytics) (11.0.2.54)
Requirement already satisfied: nvidia-curand-cu12==10.3.2.106 in
/usr/local/lib/python3.10/dist-packages (from torch>=1.8.0-
>ultralytics) (10.3.2.106)
Requirement already satisfied: nvidia-cusolver-cu12==11.4.5.107 in
/usr/local/lib/python3.10/dist-packages (from torch>=1.8.0-
>ultralytics) (11.4.5.107)
Requirement already satisfied: nvidia-cusparse-cu12==12.1.0.106 in
/usr/local/lib/python3.10/dist-packages (from torch>=1.8.0-
>ultralytics) (12.1.0.106)
Requirement already satisfied: nvidia-nccl-cu12==2.20.5 in
/usr/local/lib/python3.10/dist-packages (from torch>=1.8.0-
>ultralytics) (2.20.5)
Requirement already satisfied: nvidia-nvtx-cu12==12.1.105 in
/usr/local/lib/python3.10/dist-packages (from torch>=1.8.0-
>ultralytics) (12.1.105)
Requirement already satisfied: triton==2.3.0 in
/usr/local/lib/python3.10/dist-packages (from torch>=1.8.0-
>ultralytics) (2.3.0)
Requirement already satisfied: nvidia-nvjitlink-cu12 in
/usr/local/lib/python3.10/dist-packages (from nvidia-cusolver-
cu12==11.4.5.107->torch>=1.8.0->ultralytics) (12.5.40)
Requirement already satisfied: six>=1.5 in
/usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7-
>matplotlib>=3.3.0->ultralytics) (1.16.0)
Requirement already satisfied: MarkupSafe>=2.0 in
/usr/local/lib/python3.10/dist-packages (from jinja2->torch>=1.8.0-
>ultralytics) (2.1.5)
Requirement already satisfied: mpmath<1.4.0,>=1.1.0 in
/usr/local/lib/python3.10/dist-packages (from sympy->torch>=1.8.0-
>ultralytics) (1.3.0)
Requirement already satisfied: matplotlib in
/usr/local/lib/python3.10/dist-packages (3.7.1)
Requirement already satisfied: contourpy>=1.0.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (1.2.1)
Requirement already satisfied: cycler>=0.10 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (4.53.0)
```

```
Requirement already satisfied: kiwisolver>=1.0.1 in  
/usr/local/lib/python3.10/dist-packages (from matplotlib) (1.4.5)  
Requirement already satisfied: numpy>=1.20 in  
/usr/local/lib/python3.10/dist-packages (from matplotlib) (1.25.2)  
Requirement already satisfied: packaging>=20.0 in  
/usr/local/lib/python3.10/dist-packages (from matplotlib) (24.1)  
Requirement already satisfied: pillow>=6.2.0 in  
/usr/local/lib/python3.10/dist-packages (from matplotlib) (9.4.0)  
Requirement already satisfied: pyparsing>=2.3.1 in  
/usr/local/lib/python3.10/dist-packages (from matplotlib) (3.1.2)  
Requirement already satisfied: python-dateutil>=2.7 in  
/usr/local/lib/python3.10/dist-packages (from matplotlib) (2.8.2)  
Requirement already satisfied: six>=1.5 in  
/usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7-  
>matplotlib) (1.16.0)  
Requirement already satisfied: opencv-python-headless in  
/usr/local/lib/python3.10/dist-packages (4.10.0.82)  
Requirement already satisfied: numpy>=1.21.2 in  
/usr/local/lib/python3.10/dist-packages (from opencv-python-headless)  
(1.25.2)
```

image 1/1 /content/Emue.jpg: 480x640 1 bicycle, 1 car, 1 truck, 1 dog,  
204.5ms

Speed: 4.0ms preprocess, 204.5ms inference, 1.2ms postprocess per  
image at shape (1, 3, 480, 640)

